



Deep Learning for Poets (Part II)

Amir H. Payberah
payberah@kth.se
19/12/2018





TensorFlow

Linear and Logistic
regression

Deep Feedforward
Networks

CNN, RNN, Autoencoders



TensorFlow

Linear and Logistic
regression

Deep Feedforward
Networks

CNN, RNN, Autoencoders



Linear Algebra Review

Vector

- A **vector** is an **array of numbers**.

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

Vector

- ▶ A **vector** is an **array of numbers**.
- ▶ Notation:
 - Denoted by **bold lowercase letters**, e.g., **x**.
 - **x_i** denotes the **i**th entry.

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$



Matrix and Tensor

- A **matrix** is a 2-D array of numbers.

$$\mathbf{A} = \begin{bmatrix} a_{1,1} & a_{1,2} & a_{1,3} & \dots & a_{1,n} \\ a_{2,1} & a_{2,2} & a_{2,3} & \dots & a_{2,n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{m,1} & a_{m,2} & a_{m,3} & \dots & a_{m,n} \end{bmatrix}$$



Matrix and Tensor

- ▶ A **matrix** is a 2-D array of numbers.
- ▶ A **tensor** is an array with more than two axes.

$$\mathbf{A} = \begin{bmatrix} a_{1,1} & a_{1,2} & a_{1,3} & \dots & a_{1,n} \\ a_{2,1} & a_{2,2} & a_{2,3} & \dots & a_{2,n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{m,1} & a_{m,2} & a_{m,3} & \dots & a_{m,n} \end{bmatrix}$$



Matrix and Tensor

- ▶ A **matrix** is a 2-D array of numbers.
- ▶ A **tensor** is an array with more than two axes.
- ▶ Notation:
 - Denoted by **bold uppercase letters**, e.g., **A**.
 - a_{ij} denotes the entry in *i*th row and *j*th column.
 - If **A** is $m \times n$, it has *m* rows and *n* columns.

$$\mathbf{A} = \begin{bmatrix} a_{1,1} & a_{1,2} & a_{1,3} & \dots & a_{1,n} \\ a_{2,1} & a_{2,2} & a_{2,3} & \dots & a_{2,n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{m,1} & a_{m,2} & a_{m,3} & \dots & a_{m,n} \end{bmatrix}$$




Matrix Addition and Subtraction

- The **matrices** must have the **same dimensions**.

$$\mathbf{A} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} + \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} a+e & b+f \\ c+g & d+h \end{bmatrix}$$

Matrix Product

- The matrix product: $\mathbf{C} = \mathbf{AB}$.




[https://en.wikipedia.org/wiki/Matrix_multiplication]

Matrix Product

- ▶ The matrix product: $\mathbf{C} = \mathbf{AB}$.
- ▶ If \mathbf{A} is of shape $m \times n$ and \mathbf{B} is of shape $n \times p$, then \mathbf{C} is of shape $m \times p$.

$$c_{ij} = \sum_k a_{ik} b_{kj}$$




[https://en.wikipedia.org/wiki/Matrix_multiplication]

Matrix Product

- ▶ The matrix product: $\mathbf{C} = \mathbf{AB}$.
- ▶ If \mathbf{A} is of shape $m \times n$ and \mathbf{B} is of shape $n \times p$, then \mathbf{C} is of shape $m \times p$.

$$c_{ij} = \sum_k a_{ik} b_{kj}$$

- ▶ Properties
 - Associative: $(\mathbf{AB})\mathbf{C} = \mathbf{A}(\mathbf{BC})$
 - Not commutative: $\mathbf{AB} \neq \mathbf{BA}$



[https://en.wikipedia.org/wiki/Matrix_multiplication]



Matrix Transpose

- ▶ Swap the rows and columns of a matrix.

$$\mathbf{A} = \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \Rightarrow \mathbf{A}^T = \begin{bmatrix} a & c & e \\ b & d & f \end{bmatrix}$$



Matrix Transpose

- ▶ Swap the rows and columns of a matrix.

$$\mathbf{A} = \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \Rightarrow \mathbf{A}^T = \begin{bmatrix} a & c & e \\ b & d & f \end{bmatrix}$$

- ▶ Properties

- $\mathbf{A}_{ij} = \mathbf{A}_{ji}^T$
- If \mathbf{A} is $m \times n$, then \mathbf{A}^T is $n \times m$
- $(\mathbf{A} + \mathbf{B})^T = \mathbf{A}^T + \mathbf{B}^T$
- $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$



Inverse of a Matrix

- If \mathbf{A} is a **square** matrix, its **inverse** is called \mathbf{A}^{-1} .

$$\mathbf{A}\mathbf{A}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$$



Inverse of a Matrix

- If \mathbf{A} is a **square** matrix, its **inverse** is called \mathbf{A}^{-1} .

$$\mathbf{A}\mathbf{A}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$$

- Where \mathbf{I} , the **identity** matrix, is a **diagonal matrix** with all **1's** on the diagonal.

$$\mathbf{I}_2 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad \mathbf{I}_3 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



L^p Norm for Vectors

- We can measure the **size of vectors** using a **norm** function.



L^p Norm for Vectors

- ▶ We can measure the **size of vectors** using a **norm** function.
- ▶ Norms are functions **mapping vectors to non-negative values**.



L^p Norm for Vectors

- ▶ We can measure the **size of vectors** using a **norm** function.
- ▶ Norms are functions **mapping vectors to non-negative values**.
- ▶ L¹ norm

$$\|x\|_1 = \sum_i |x_i|$$



L^p Norm for Vectors

- ▶ We can measure the **size of vectors** using a **norm** function.
- ▶ Norms are functions **mapping vectors to non-negative values**.
- ▶ L¹ norm

$$\|\mathbf{x}\|_1 = \sum_i |x_i|$$

- ▶ L² norm

$$\|\mathbf{x}\|_2 = \left(\sum_i |x_i|^2 \right)^{\frac{1}{2}} = \sqrt{x_1^2 + x_2^2 + \dots + x_n^2}$$



L^p Norm for Vectors

- We can measure the **size of vectors** using a **norm** function.
- Norms are functions **mapping vectors to non-negative values**.
- L¹ norm


$$\|\mathbf{x}\|_1 = \sum_i |x_i|$$

- L² norm

$$\|\mathbf{x}\|_2 = \left(\sum_i |x_i|^2 \right)^{\frac{1}{2}} = \sqrt{x_1^2 + x_2^2 + \dots + x_n^2}$$

- L^p norm

$$\|\mathbf{x}\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$$





Probability Review



Random Variables

- ▶ Random variable: a **variable** that can take on **different values randomly**.



Random Variables

- ▶ **Random variable:** a **variable** that can take on **different values randomly**.
- ▶ Random variables may be **discrete** or **continuous**.



Random Variables

- ▶ **Random variable:** a **variable** that can take on **different values randomly**.
- ▶ Random variables may be **discrete** or **continuous**.
- ▶ **Notation:**
 - Denoted by an **upper case letter**, e.g., **X**
 - Values of a random variable **X** are denoted by **lower case letters**, e.g., **x** and **y**.



Probability Distributions

- ▶ **Probability distribution:** how likely a **random variable** is to take on each of its **possible states**.



Probability Distributions

- ▶ **Probability distribution:** how likely a **random variable** is to take on each of its **possible states**.
- ▶ E.g., the **random variable** **X** denotes the outcome of a **coin toss**.



Probability Distributions

- ▶ **Probability distribution:** how likely a **random variable** is to take on each of its **possible states**.
- ▶ E.g., the **random variable** **X** denotes the outcome of a **coin toss**.
 - The **probability distribution** of **X** would take the value **0.5** for **X = head**, and **0.5** for **Y = tail** (assuming the coin is **fair**).



Probability Distributions

- ▶ **Probability distribution:** how likely a **random variable** is to take on each of its **possible states**.
- ▶ E.g., the **random variable** **X** denotes the outcome of a **coin toss**.
 - The **probability distribution** of **X** would take the value **0.5** for **X = head**, and **0.5** for **Y = tail** (assuming the coin is **fair**).
- ▶ The way we **describe probability distributions** depends on whether the variables are **discrete** or **continuous**.



Discrete Variables

- ▶ Probability mass function (PMF): the probability distribution of a discrete random variable X .



Discrete Variables

- ▶ Probability mass function (PMF): the probability distribution of a discrete random variable X .
- ▶ Notation: denoted by a lowercase p .
 - E.g., $p(x) = 1$ indicates that $X = x$ is certain
 - E.g., $p(x) = 0$ indicates that $X = x$ is impossible



Discrete Variables

- ▶ Probability mass function (PMF): the probability distribution of a discrete random variable X .
- ▶ Notation: denoted by a lowercase p .
 - E.g., $p(x) = 1$ indicates that $X = x$ is certain
 - E.g., $p(x) = 0$ indicates that $X = x$ is impossible
- ▶ Properties:
 - The domain D of p must be the set of all possible states of X
 - $\forall x \in D(X), 0 \leq p(x) \leq 1$
 - $\sum_{x \in D(X)} p(x) = 1$



Independence

- ▶ Two random variables X and Y are **independent**, if their **probability distribution** can be expressed as their **products**.

$$\forall x \in D(X), y \in D(Y), p(X = x, Y = y) = p(X = x)p(Y = y)$$



Independence

- ▶ Two random variables X and Y are **independent**, if their **probability distribution** can be expressed as their **products**.

$$\forall x \in D(X), y \in D(Y), p(X = x, Y = y) = p(X = x)p(Y = y)$$

- ▶ E.g., if a **coin is tossed** and a single **6-sided die is rolled**, then the probability of landing on the **head** side of the coin and **rolling a 3** on the die is:

Independence

- ▶ Two random variables X and Y are **independent**, if their **probability distribution** can be expressed as their **products**.

$$\forall x \in D(X), y \in D(Y), p(X = x, Y = y) = p(X = x)p(Y = y)$$

- ▶ E.g., if a **coin is tossed** and a single **6-sided die is rolled**, then the probability of landing on the **head** side of the coin and **rolling a 3** on the die is:

$$p(X = \text{head}, Y = 3) = p(X = \text{head})p(Y = 3) = \frac{1}{2} \times \frac{1}{6} = \frac{1}{12}$$



Conditional Probability

- ▶ **Conditional probability:** the probability of an event given that another event has occurred.

$$p(Y = y \mid X = x) = \frac{p(Y = y, X = x)}{p(X = x)}$$



Conditional Probability

- ▶ **Conditional probability:** the probability of an event given that another event has occurred.

$$p(Y = y \mid X = x) = \frac{p(Y = y, X = x)}{p(X = x)}$$

- ▶ E.g., if 60% of the class passed both labs and 80% of the class passed the first labs, then what percent of those who passed the first lab also passed the second lab?



Conditional Probability

- ▶ **Conditional probability:** the probability of an event given that another event has occurred.

$$p(Y = y \mid X = x) = \frac{p(Y = y, X = x)}{p(X = x)}$$

- ▶ E.g., if 60% of the class passed both labs and 80% of the class passed the first labs, then what percent of those who passed the first lab also passed the second lab?
 - E.g., X and Y random variables for the first and the second labs, respectively.

$$p(Y = \text{lab2} \mid X = \text{lab1}) = \frac{p(Y = \text{lab2}, X = \text{lab1})}{p(X = \text{lab1})} = \frac{0.6}{0.8} = \frac{3}{4}$$



Expectation

- ▶ The **expected value** of a random variable X with respect to a probability distribution $p(x)$ is the **average** value that X takes on when it is drawn from $p(x)$.

$$E_{x \sim p}[X] = \sum_x p(x)x$$

Expectation

- ▶ The **expected value** of a random variable X with respect to a probability distribution $p(x)$ is the **average** value that X takes on when it is drawn from $p(x)$.

$$E_{x \sim p}[X] = \sum_x p(x)x$$

- ▶ E.g., If $X : \{1, 2, 3\}$, and $p(X = 1) = 0.3$, $p(X = 2) = 0.5$, $p(X = 3) = 0.2$

Expectation

- ▶ The **expected value** of a random variable X with respect to a probability distribution $p(x)$ is the **average** value that X takes on when it is drawn from $p(x)$.

$$E_{x \sim p}[X] = \sum_x p(x)x$$

- ▶ E.g., If $X : \{1, 2, 3\}$, and $p(X = 1) = 0.3$, $p(X = 2) = 0.5$, $p(X = 3) = 0.2$
 - $E[X] = 0.3 \times 1 + 0.5 \times 2 + 0.2 \times 3 = 1.9$



Variance and Standard Deviation

- ▶ The **variance** gives a measure of how much the **values** of a random variable **X** vary as we sample it from its **probability distribution p(X)**.

$$\text{Var}(X) = E[(X - E[X])^2]$$

$$\text{Var}(X) = \sum_x p(x)(x - E[X])^2$$



Variance and Standard Deviation

- ▶ The **variance** gives a measure of how much the **values** of a random variable **X** vary as we sample it from its **probability distribution p(X)**.

$$\text{Var}(X) = E[(X - E[X])^2]$$

$$\text{Var}(X) = \sum_x p(x)(x - E[X])^2$$

- ▶ E.g., If $X : \{1, 2, 3\}$, and $p(X = 1) = 0.3$, $p(X = 2) = 0.5$, $p(X = 3) = 0.2$



Variance and Standard Deviation

- The **variance** gives a measure of how much the **values** of a random variable **X** vary as we sample it from its **probability distribution p(X)**.

$$\text{Var}(X) = E[(X - E[X])^2]$$

$$\text{Var}(X) = \sum_x p(x)(x - E[X])^2$$

- E.g., If $X : \{1, 2, 3\}$, and $p(X = 1) = 0.3$, $p(X = 2) = 0.5$, $p(X = 3) = 0.2$
 - $E[X] = 0.3 \times 1 + 0.5 \times 2 + 0.2 \times 3 = 1.9$
 - $\text{Var}(X) = 0.3(1 - 1.9)^2 + 0.5(2 - 1.9)^2 + 0.2(3 - 1.9)^2 = 0.49$



Variance and Standard Deviation

- ▶ The **variance** gives a measure of how much the **values** of a random variable **X** vary as we sample it from its **probability distribution p(X)**.

$$\text{Var}(X) = E[(X - E[X])^2]$$

$$\text{Var}(X) = \sum_x p(x)(x - E[X])^2$$

- ▶ E.g., If $X : \{1, 2, 3\}$, and $p(X = 1) = 0.3$, $p(X = 2) = 0.5$, $p(X = 3) = 0.2$
 - $E[X] = 0.3 \times 1 + 0.5 \times 2 + 0.2 \times 3 = 1.9$
 - $\text{Var}(X) = 0.3(1 - 1.9)^2 + 0.5(2 - 1.9)^2 + 0.2(3 - 1.9)^2 = 0.49$
- ▶ The **standard deviation**, shown by σ , is the **square root of the variance**.



Probability and Likelihood (1/2)

- ▶ Let $X : \{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ be a discrete random variable drawn independently from a distribution probability p depending on a parameter θ .



Probability and Likelihood (1/2)

- ▶ Let $X : \{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ be a discrete random variable drawn independently from a distribution probability p depending on a parameter θ .
 - For six tosses of a coin, $X : \{h, t, t, t, h, t\}$, h : head, and t : tail.
 - Suppose you have a coin with probability θ to land heads and $(1 - \theta)$ to land tails.



Probability and Likelihood (1/2)

- ▶ Let $X : \{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ be a discrete random variable drawn independently from a distribution probability p depending on a parameter θ .
 - For six tosses of a coin, $X : \{h, t, t, t, h, t\}$, h : head, and t : tail.
 - Suppose you have a coin with probability θ to land heads and $(1 - \theta)$ to land tails.
- ▶ $p(X | \theta = \frac{2}{3})$ is the probability of X given $\theta = \frac{2}{3}$.



Probability and Likelihood (1/2)

- ▶ Let $X : \{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ be a **discrete random variable** drawn independently from a **distribution probability p** depending on a **parameter θ** .
 - For six tosses of a coin, $X : \{h, t, t, t, h, t\}$, **h**: head, and **t**: tail.
 - Suppose you have a **coin** with probability θ to land heads and $(1 - \theta)$ to land tails.
- ▶ $p(X | \theta = \frac{2}{3})$ is the **probability** of X given $\theta = \frac{2}{3}$.
- ▶ $p(X = h | \theta)$ is the **likelihood** of θ given $X = h$.

Probability and Likelihood (1/2)

- ▶ Let $X : \{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ be a discrete random variable drawn independently from a distribution probability p depending on a parameter θ .
 - For six tosses of a coin, $X : \{h, t, t, t, h, t\}$, h : head, and t : tail.
 - Suppose you have a coin with probability θ to land heads and $(1 - \theta)$ to land tails.
- ▶ $p(X | \theta = \frac{2}{3})$ is the probability of X given $\theta = \frac{2}{3}$.
- ▶ $p(X = h | \theta)$ is the likelihood of θ given $X = h$.
- ▶ Likelihood (L): a function of the parameters (θ) of a probability model, given specific observed data, e.g., $X = h$.

$$L(\theta | X) = p(X | \theta)$$



Probability and Likelihood (2/2)

- ▶ The **likelihood** differs from that of a **probability**.



Probability and Likelihood (2/2)

- ▶ The **likelihood** differs from that of a **probability**.
- ▶ A probability $p(X | \theta)$ refers to the occurrence of **future events**.



Probability and Likelihood (2/2)

- ▶ The likelihood differs from that of a probability.
- ▶ A probability $p(X | \theta)$ refers to the occurrence of future events.
- ▶ A likelihood $L(\theta | X)$ refers to past events with known outcomes.



Likelihood and Log-Likelihood (1/2)

- If samples in X are **independent** we have:

$$\begin{aligned} L(\theta \mid X) &= p(X \mid \theta) = p(x^{(1)}, x^{(2)}, \dots, x^{(m)} \mid \theta) \\ &= p(x^{(1)} \mid \theta)p(x^{(2)} \mid \theta) \cdots p(x^{(m)} \mid \theta) = \prod_{i=1}^m p(x^{(i)} \mid \theta) \end{aligned}$$

Likelihood and Log-Likelihood (1/2)

- If samples in X are **independent** we have:

$$\begin{aligned} L(\theta \mid X) &= p(X \mid \theta) = p(x^{(1)}, x^{(2)}, \dots, x^{(m)} \mid \theta) \\ &= p(x^{(1)} \mid \theta)p(x^{(2)} \mid \theta) \cdots p(x^{(m)} \mid \theta) = \prod_{i=1}^m p(x^{(i)} \mid \theta) \end{aligned}$$

- E.g., **six tosses** of a coin, with the following model:

- Data: $X : \{h, t, t, t, h, t\}$
- Possible **outcomes**: h with probability of θ , and t with probability $(1 - \theta)$.

Likelihood and Log-Likelihood (1/2)

- If samples in X are **independent** we have:

$$\begin{aligned} L(\theta \mid X) &= p(X \mid \theta) = p(x^{(1)}, x^{(2)}, \dots, x^{(m)} \mid \theta) \\ &= p(x^{(1)} \mid \theta)p(x^{(2)} \mid \theta) \cdots p(x^{(m)} \mid \theta) = \prod_{i=1}^m p(x^{(i)} \mid \theta) \end{aligned}$$

- E.g., **six tosses** of a coin, with the following model:

- Data: $X : \{h, t, t, t, h, t\}$
- Possible **outcomes**: h with probability of θ , and t with probability $(1 - \theta)$.

$$\begin{aligned} L(\theta \mid X) &= p(X \mid \theta) \\ &= p(X = h \mid \theta)p(X = t \mid \theta)p(X = t \mid \theta)p(X = t \mid \theta)p(X = h \mid \theta)p(X = t \mid \theta) \\ &= \theta(1 - \theta)(1 - \theta)(1 - \theta)\theta(1 - \theta) \\ &= \theta^2(1 - \theta)^4 \end{aligned}$$

Likelihood and Log-Likelihood (2/2)

- ▶ The probability product is prone to numerical underflow.

$$L(\theta \mid X) = p(X \mid \theta) = \prod_{i=1}^m p(x^{(i)} \mid \theta)$$

Likelihood and Log-Likelihood (2/2)

- ▶ The probability product is prone to numerical underflow.

$$L(\theta \mid X) = p(X \mid \theta) = \prod_{i=1}^m p(x^{(i)} \mid \theta)$$

- ▶ To overcome this problem we can use the logarithm of the likelihood.

$$\log L(\theta \mid X) = \log \prod_{i=1}^m p(x^{(i)} \mid \theta) = \sum_{i=1}^m \log p(x^{(i)} \mid \theta)$$



Negative Log-Likelihood

- ▶ Likelihood: $L(\theta \mid X) = \prod_{i=1}^m p(x^{(i)} \mid \theta)$



Negative Log-Likelihood

- ▶ Likelihood: $L(\theta | X) = \prod_{i=1}^m p(x^{(i)} | \theta)$
- ▶ Log-Likelihood: $\log L(\theta | X) = \log \prod_{i=1}^m p(x^{(i)} | \theta) = \sum_{i=1}^m \log p(x^{(i)} | \theta)$



Negative Log-Likelihood

- ▶ Likelihood: $L(\theta | X) = \prod_{i=1}^m p(x^{(i)} | \theta)$
- ▶ Log-Likelihood: $\log L(\theta | X) = \log \prod_{i=1}^m p(x^{(i)} | \theta) = \sum_{i=1}^m \log p(x^{(i)} | \theta)$
- ▶ Negative Log-Likelihood: $-\log L(\theta | X) = -\sum_{i=1}^m \log p(x^{(i)} | \theta)$



Negative Log-Likelihood

- ▶ Likelihood: $L(\theta | X) = \prod_{i=1}^m p(x^{(i)} | \theta)$
- ▶ Log-Likelihood: $\log L(\theta | X) = \log \prod_{i=1}^m p(x^{(i)} | \theta) = \sum_{i=1}^m \log p(x^{(i)} | \theta)$
- ▶ Negative Log-Likelihood: $-\log L(\theta | X) = -\sum_{i=1}^m \log p(x^{(i)} | \theta)$
- ▶ Negative log-likelihood is also called the **cross-entropy**



Cross-Entropy

- ▶ **Cross-entropy**: quantify the **difference (error)** between **two probability distributions**.
- ▶ How close is the **predicted distribution** to the **true distribution**?

$$H(p, q) = - \sum_x p(x) \log(q(x))$$

- ▶ Where **p** is the **true distribution**, and **q** the **predicted distribution**.



Cross-Entropy - Example

- ▶ Six tosses of a coin: $X : \{h, t, t, t, h, t\}$
- ▶ The true distribution p : $p(h) = \frac{2}{6}$ and $p(t) = \frac{4}{6}$
- ▶ The predicted distribution q : h with probability of θ , and t with probability $(1 - \theta)$.



Cross-Entropy - Example

- ▶ Six tosses of a coin: $X : \{h, t, t, t, h, t\}$
- ▶ The **true distribution** p : $p(h) = \frac{2}{6}$ and $p(t) = \frac{4}{6}$
- ▶ The **predicted distribution** q : h with probability of θ , and t with probability $(1 - \theta)$.
- ▶ Cross entropy: $H(p, q) = -\sum_x p(x)\log(q(x))$
 $= -p(h)\log(q(h)) - p(t)\log(q(t)) = -\frac{2}{6}\log(\theta) - \frac{4}{6}\log(1 - \theta)$



Cross-Entropy - Example

- ▶ Six tosses of a coin: $X : \{h, t, t, t, h, t\}$
- ▶ The **true distribution** p : $p(h) = \frac{2}{6}$ and $p(t) = \frac{4}{6}$
- ▶ The **predicted distribution** q : h with probability of θ , and t with probability $(1 - \theta)$.
- ▶ Cross entropy: $H(p, q) = -\sum_x p(x)\log(q(x))$
 $= -p(h)\log(q(h)) - p(t)\log(q(t)) = -\frac{2}{6}\log(\theta) - \frac{4}{6}\log(1 - \theta)$
- ▶ Likelihood: $\theta^2(1 - \theta)^4$



Cross-Entropy - Example

- ▶ Six tosses of a coin: $X : \{h, t, t, t, h, t\}$
- ▶ The **true distribution** p : $p(h) = \frac{2}{6}$ and $p(t) = \frac{4}{6}$
- ▶ The **predicted distribution** q : h with probability of θ , and t with probability $(1 - \theta)$.
- ▶ Cross entropy: $H(p, q) = -\sum_x p(x)\log(q(x))$
 $= -p(h)\log(q(h)) - p(t)\log(q(t)) = -\frac{2}{6}\log(\theta) - \frac{4}{6}\log(1 - \theta)$
- ▶ Likelihood: $\theta^2(1 - \theta)^4$
- ▶ Negative log likelihood: $-\log(\theta^2(1 - \theta)^4) = -2\log(\theta) - 4\log(1 - \theta)$



Linear Regression



Let's Start with an Example



The Housing Price Example (1/3)

- ▶ Given the dataset of m houses.

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |



The Housing Price Example (1/3)

- Given the dataset of m houses.

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |

- Predict the prices of other houses, as a function of the size of living area and number of bedrooms?



The Housing Price Example (2/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |
| : | : | : |

The Housing Price Example (2/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |

$$\mathbf{x}^{(1)} = \begin{bmatrix} 2104 \\ 3 \end{bmatrix} \quad y^{(1)} = 400 \quad \mathbf{x}^{(2)} = \begin{bmatrix} 1600 \\ 3 \end{bmatrix} \quad y^{(2)} = 330 \quad \mathbf{x}^{(3)} = \begin{bmatrix} 2400 \\ 3 \end{bmatrix} \quad y^{(3)} = 369$$

The Housing Price Example (2/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |

$$\mathbf{x}^{(1)} = \begin{bmatrix} 2104 \\ 3 \end{bmatrix} \quad y^{(1)} = 400 \quad \mathbf{x}^{(2)} = \begin{bmatrix} 1600 \\ 3 \end{bmatrix} \quad y^{(2)} = 330 \quad \mathbf{x}^{(3)} = \begin{bmatrix} 2400 \\ 3 \end{bmatrix} \quad y^{(3)} = 369$$

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}^{(1)\top} \\ \mathbf{x}^{(2)\top} \\ \mathbf{x}^{(3)\top} \\ \vdots \end{bmatrix} = \begin{bmatrix} 2104 & 3 \\ 1600 & 3 \\ 2400 & 3 \\ \vdots & \vdots \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 400 \\ 330 \\ 369 \\ \vdots \end{bmatrix}$$

The Housing Price Example (2/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |

$$\mathbf{x}^{(1)} = \begin{bmatrix} 2104 \\ 3 \end{bmatrix} \quad y^{(1)} = 400 \quad \mathbf{x}^{(2)} = \begin{bmatrix} 1600 \\ 3 \end{bmatrix} \quad y^{(2)} = 330 \quad \mathbf{x}^{(3)} = \begin{bmatrix} 2400 \\ 3 \end{bmatrix} \quad y^{(3)} = 369$$

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}^{(1)\top} \\ \mathbf{x}^{(2)\top} \\ \mathbf{x}^{(3)\top} \\ \vdots \end{bmatrix} = \begin{bmatrix} 2104 & 3 \\ 1600 & 3 \\ 2400 & 3 \\ \vdots & \vdots \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 400 \\ 330 \\ 369 \\ \vdots \end{bmatrix}$$

- $\mathbf{x}^{(i)} \in \mathbb{R}^2$: $x_1^{(i)}$ is the living area, and $x_2^{(i)}$ is the number of bedrooms of the i th house in the training set.

The Housing Price Example (3/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |

- ▶ Predict the prices of other houses \hat{y} as a function of the size of their living areas x_1 , and number of bedrooms x_2 , i.e., $\hat{y} = f(x_1, x_2)$
- ▶ E.g., what is \hat{y} , if $x_1 = 4000$ and $x_2 = 4$?

The Housing Price Example (3/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| : | : | : |

- ▶ Predict the prices of other houses \hat{y} as a function of the size of their living areas x_1 , and number of bedrooms x_2 , i.e., $\hat{y} = f(x_1, x_2)$
- ▶ E.g., what is \hat{y} , if $x_1 = 4000$ and $x_2 = 4$?
- ▶ As an initial choice: $\hat{y} = f_w(\mathbf{x}) = w_1x_1 + w_2x_2$



Linear Regression (1/2)

- ▶ Our goal: to build a system that takes input $\mathbf{x} \in \mathbb{R}^n$ and predicts output $\hat{\mathbf{y}} \in \mathbb{R}$.



Linear Regression (1/2)

- ▶ Our goal: to build a system that takes input $\mathbf{x} \in \mathbb{R}^n$ and predicts output $\hat{\mathbf{y}} \in \mathbb{R}$.
- ▶ In linear regression, the output $\hat{\mathbf{y}}$ is a linear function of the input \mathbf{x} .

$$\begin{aligned}\hat{\mathbf{y}} &= f_w(\mathbf{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n \\ \hat{\mathbf{y}} &= \mathbf{w}^\top \mathbf{x}\end{aligned}$$



Linear Regression (1/2)

- ▶ Our goal: to build a system that takes input $\mathbf{x} \in \mathbb{R}^n$ and predicts output $\hat{\mathbf{y}} \in \mathbb{R}$.
- ▶ In linear regression, the output $\hat{\mathbf{y}}$ is a linear function of the input \mathbf{x} .


$$\begin{aligned}\hat{\mathbf{y}} &= f_{\mathbf{w}}(\mathbf{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n \\ \hat{\mathbf{y}} &= \mathbf{w}^T \mathbf{x}\end{aligned}$$

- $\hat{\mathbf{y}}$: the predicted value
- x_i : the i th feature value
- w_j : the j th model parameter ($\mathbf{w} \in \mathbb{R}^n$)
- n : the number of features

Linear Regression (2/2)

- ▶ Linear regression often has one additional parameter, called **intercept b**:


$$\hat{y} = \mathbf{w}^T \mathbf{x} + b$$



Linear Regression (2/2)

- ▶ Linear regression often has one additional parameter, called **intercept b**:

$$\hat{y} = \mathbf{w}^T \mathbf{x} + b$$



- ▶ Instead of adding the bias parameter **b**, we can augment **x** with an **extra entry** that is **always set to 1**.

$$\hat{y} = f_w(\mathbf{x}) = w_0 x_0 + w_1 x_1 + w_2 x_2 + \cdots + w_n x_n, \text{ where } x_0 = 1$$



Linear Regression - Model Parameters

- ▶ Parameters $w \in \mathbb{R}^n$ are values that control the behavior of the model.



Linear Regression - Model Parameters



- ▶ Parameters $w \in \mathbb{R}^n$ are values that control the behavior of the model.
- ▶ w are a set of weights that determine how each feature affects the prediction.



How to Learn Model Parameters w ?





Linear Regression - Cost Function (1/2)



- One reasonable model should make \hat{y} close to y , at least for the training dataset.


Linear Regression - Cost Function (1/2)



- One reasonable model should make \hat{y} close to y , at least for the training dataset.
- Residual: the difference between the dependent variable y and the predicted value \hat{y} .

$$r^{(i)} = y^{(i)} - \hat{y}^{(i)}$$

Linear Regression - Cost Function (2/2)




► Cost function $J(\mathbf{w})$

- For each value of the \mathbf{w} , it measures how **close** the $\hat{y}^{(i)}$ is to the corresponding $y^{(i)}$.
- We can define $J(\mathbf{w})$ as the **mean squared error (MSE)**:

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

Linear Regression - Cost Function (2/2)



► Cost function $J(\mathbf{w})$

- For each value of the \mathbf{w} , it measures how close the $\hat{y}^{(i)}$ is to the corresponding $y^{(i)}$.
- We can define $J(\mathbf{w})$ as the mean squared error (MSE):

$$\begin{aligned} J(\mathbf{w}) &= \text{MSE}(\mathbf{w}) = \frac{1}{m} \sum_{i}^m (\hat{y}^{(i)} - y^{(i)})^2 \\ &= E[(\hat{y} - y)^2] = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 \end{aligned}$$



How to Learn Model Parameters?


- ▶ We want to choose \mathbf{w} so as to minimize $J(\mathbf{w})$.
- ▶ Two approaches to find \mathbf{w} :
 - Normal equation
 - Gradient descent



Normal Equation


Derivatives and Gradient (1/3)

- The **first derivative** of $f(x)$, shown as $f'(x)$, shows the **slope** of the **tangent line** to the function at the point x .




Derivatives and Gradient (1/3)

- ▶ The **first derivative** of $f(x)$, shown as $f'(x)$, shows the **slope** of the **tangent line** to the function at the point x .
- ▶ $f(x) = x^2 \Rightarrow f'(x) = 2x$




Derivatives and Gradient (1/3)

- ▶ The **first derivative** of $f(x)$, shown as $f'(x)$, shows the **slope** of the **tangent line** to the function at the point x .
- ▶ $f(x) = x^2 \Rightarrow f'(x) = 2x$
- ▶ If $f(x)$ is increasing, then $f'(x) > 0$




Derivatives and Gradient (1/3)

- ▶ The **first derivative** of $f(x)$, shown as $f'(x)$, shows the **slope** of the **tangent line** to the function at the point x .
- ▶ $f(x) = x^2 \Rightarrow f'(x) = 2x$
- ▶ If $f(x)$ is increasing, then $f'(x) > 0$
- ▶ If $f(x)$ is decreasing, then $f'(x) < 0$



Derivatives and Gradient (1/3)

- ▶ The **first derivative** of $f(x)$, shown as $f'(x)$, shows the **slope** of the **tangent line** to the function at the point x .
- ▶ $f(x) = x^2 \Rightarrow f'(x) = 2x$
- ▶ If $f(x)$ is increasing, then $f'(x) > 0$
- ▶ If $f(x)$ is decreasing, then $f'(x) < 0$
- ▶ If $f(x)$ is at local minimum/maximum, then $f'(x) = 0$





Derivatives and Gradient (2/3)

- ▶ What if a function has multiple arguments, e.g., $f(x_1, x_2, \dots, x_n)$



Derivatives and Gradient (2/3)

- ▶ What if a function has multiple arguments, e.g., $f(x_1, x_2, \dots, x_n)$
- ▶ **Partial derivatives:** the derivative with respect to a **particular argument**.
 - $\frac{\partial f}{\partial x_1}$, the derivative **with respect to x_1**
 - $\frac{\partial f}{\partial x_2}$, the derivative **with respect to x_2**



Derivatives and Gradient (2/3)

- ▶ What if a function has multiple arguments, e.g., $f(x_1, x_2, \dots, x_n)$
- ▶ **Partial derivatives:** the derivative with respect to a particular argument.
 - $\frac{\partial f}{\partial x_1}$, the derivative with respect to x_1
 - $\frac{\partial f}{\partial x_2}$, the derivative with respect to x_2
- ▶ $\frac{\partial f}{\partial x_i}$: shows how much the function f will change, if we change x_i .

Derivatives and Gradient (2/3)

- ▶ What if a function has multiple arguments, e.g., $f(x_1, x_2, \dots, x_n)$
- ▶ **Partial derivatives:** the derivative with respect to a particular argument.
 - $\frac{\partial f}{\partial x_1}$, the derivative with respect to x_1
 - $\frac{\partial f}{\partial x_2}$, the derivative with respect to x_2
- ▶ $\frac{\partial f}{\partial x_i}$: shows how much the function f will change, if we change x_i .
- ▶ **Gradient:** the vector of all partial derivatives for a function f .

$$\nabla_x f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$



Derivatives and Gradient (3/3)

- ▶ What is the gradient of $f(x_1, x_2, x_3) = x_1 - x_1x_2 + x_3^2$?

Derivatives and Gradient (3/3)

- ▶ What is the gradient of $f(x_1, x_2, x_3) = x_1 - x_1x_2 + x_3^2$?

$$\nabla_x f(x) = \begin{bmatrix} \frac{\partial}{\partial x_1}(x_1 - x_1x_2 + x_3^2) \\ \frac{\partial}{\partial x_2}(x_1 - x_1x_2 + x_3^2) \\ \frac{\partial}{\partial x_3}(x_1 - x_1x_2 + x_3^2) \end{bmatrix} = \begin{bmatrix} 1 - x_2 \\ -x_1 \\ 2x_3 \end{bmatrix}$$



Normal Equation (1/2)

- ▶ To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$\hat{y} = \mathbf{w}^T \mathbf{x}$$

Normal Equation (1/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$\hat{\mathbf{y}} = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{X} = \begin{bmatrix} [x_1^{(1)}, x_2^{(1)}, \dots, x_n^{(1)}] \\ [x_1^{(2)}, x_2^{(2)}, \dots, x_n^{(2)}] \\ \vdots \\ [x_1^{(m)}, x_2^{(m)}, \dots, x_n^{(m)}] \end{bmatrix} = \begin{bmatrix} \mathbf{x}^{(1)\top} \\ \mathbf{x}^{(2)\top} \\ \vdots \\ \mathbf{x}^{(m)\top} \end{bmatrix} \quad \hat{\mathbf{y}} = \begin{bmatrix} \hat{y}^{(1)} \\ \hat{y}^{(2)} \\ \vdots \\ \hat{y}^{(m)} \end{bmatrix}$$

Normal Equation (1/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$\hat{\mathbf{y}} = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{X} = \begin{bmatrix} [x_1^{(1)}, x_2^{(1)}, \dots, x_n^{(1)}] \\ [x_1^{(2)}, x_2^{(2)}, \dots, x_n^{(2)}] \\ \vdots \\ [x_1^{(m)}, x_2^{(m)}, \dots, x_n^{(m)}] \end{bmatrix} = \begin{bmatrix} \mathbf{x}^{(1)\top} \\ \mathbf{x}^{(2)\top} \\ \vdots \\ \mathbf{x}^{(m)\top} \end{bmatrix} \quad \hat{\mathbf{y}} = \begin{bmatrix} \hat{y}^{(1)} \\ \hat{y}^{(2)} \\ \vdots \\ \hat{y}^{(m)} \end{bmatrix}$$

$$\hat{\mathbf{y}} = \mathbf{w}^T \mathbf{X}^\top \text{ or } \hat{\mathbf{y}} = \mathbf{X} \mathbf{w}$$



Normal Equation (2/2)

- ▶ To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$



Normal Equation (2/2)

- ▶ To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$J(\mathbf{w}) = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

Normal Equation (2/2)

- ▶ To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$\begin{aligned}J(\mathbf{w}) &= \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \quad \nabla_{\mathbf{w}} J(\mathbf{w}) = 0 \\&\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = 0\end{aligned}$$

Normal Equation (2/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$J(\mathbf{w}) = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = 0$$

Normal Equation (2/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$J(\mathbf{w}) = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

Normal Equation (2/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$J(\mathbf{w}) = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) = 0$$

Normal Equation (2/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$J(\mathbf{w}) = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) = 0$$

$$\Rightarrow 2\mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{X}^\top \mathbf{y} = 0$$

Normal Equation (2/2)

- To minimize $J(\mathbf{w})$, we can simply solve for where its gradient is 0: $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$

$$J(\mathbf{w}) = \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2, \nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

$$\Rightarrow \nabla_{\mathbf{w}} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) = 0$$

$$\Rightarrow 2\mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{X}^\top \mathbf{y} = 0$$

$$\Rightarrow \mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$



Normal Equation - Example (1/4)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| 1416 | 2 | 232 |
| 3000 | 4 | 540 |

- ▶ Predict the value of \hat{y} , when $x_1 = 4000$ and $x_2 = 4$.

Normal Equation - Example (1/4)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| 1416 | 2 | 232 |
| 3000 | 4 | 540 |

- ▶ Predict the value of \hat{y} , when $x_1 = 4000$ and $x_2 = 4$.
- ▶ We should find w_0 , w_1 , and w_2 in $\hat{y} = w_0 + w_1x_1 + w_2x_2$.
- ▶ $w = (X^T X)^{-1} X^T y$.

Normal Equation - Example (2/4)

| Living area | No. | of bedrooms | Price |
|-------------|-----|-------------|-------|
| 2104 | | 3 | 400 |
| 1600 | | 3 | 330 |
| 2400 | | 3 | 369 |
| 1416 | | 2 | 232 |
| 3000 | | 4 | 540 |

$$\mathbf{X} = \left[\begin{array}{ccc|c} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$



Normal Equation - Example (3/4)

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Normal Equation - Example (3/4)

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

$$\mathbf{X}^\top \mathbf{X} = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 2104 & 1600 & 2400 & 1416 & 3000 \\ 3 & 3 & 3 & 2 & 4 \end{bmatrix} \begin{bmatrix} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{bmatrix} = \begin{bmatrix} 5 & 10520 & 10520 \\ 15 & 23751872 & 33144 \\ 15 & 33144 & 47 \end{bmatrix}$$

Normal Equation - Example (3/4)

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

$$\mathbf{X}^\top \mathbf{X} = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 2104 & 1600 & 2400 & 1416 & 3000 \\ 3 & 3 & 3 & 2 & 4 \end{bmatrix} \begin{bmatrix} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{bmatrix} = \begin{bmatrix} 5 & 10520 & 15 \\ 10520 & 23751872 & 33144 \\ 15 & 33144 & 47 \end{bmatrix}$$

$$(\mathbf{X}^\top \mathbf{X})^{-1} = \begin{bmatrix} 4.90366455e+00 & 7.48766737e-04 & -2.09302326e+00 \\ 7.48766737e-04 & 2.75281889e-06 & -2.18023256e-03 \\ -2.09302326e+00 & -2.18023256e-03 & 2.22674419e+00 \end{bmatrix}$$

Normal Equation - Example (3/4)

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$\mathbf{X}^T \mathbf{X} = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 2104 & 1600 & 2400 & 1416 & 3000 \\ 3 & 3 & 3 & 2 & 4 \end{bmatrix} \begin{bmatrix} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{bmatrix} = \begin{bmatrix} 5 & 10520 & 15 \\ 10520 & 23751872 & 33144 \\ 15 & 33144 & 47 \end{bmatrix}$$

$$(\mathbf{X}^T \mathbf{X})^{-1} = \begin{bmatrix} 4.90366455e+00 & 7.48766737e-04 & -2.09302326e+00 \\ 7.48766737e-04 & 2.75281889e-06 & -2.18023256e-03 \\ -2.09302326e+00 & -2.18023256e-03 & 2.22674419e+00 \end{bmatrix}$$

$$\mathbf{X}^T \mathbf{y} = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 2104 & 1600 & 2400 & 1416 & 3000 \\ 3 & 3 & 3 & 2 & 4 \end{bmatrix} \begin{bmatrix} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{bmatrix} = \begin{bmatrix} 1871 \\ 4203712 \\ 5921 \end{bmatrix}$$

Normal Equation - Example (4/4)

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = \begin{bmatrix} 4.90366455e + 00 & 7.48766737e - 04 & -2.09302326e + 00 \\ 7.48766737e - 04 & 2.75281889e - 06 & -2.18023256e - 03 \\ -2.09302326e + 00 & -2.18023256e - 03 & 2.22674419e + 00 \end{bmatrix} \begin{bmatrix} 1871 \\ 4203712 \\ 5921 \end{bmatrix}$$
$$= \begin{bmatrix} -7.04346018e + 01 \\ 6.38433756e - 02 \\ 1.03436047e + 02 \end{bmatrix}$$



Normal Equation - Example (4/4)

$$\begin{aligned} \mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} &= \begin{bmatrix} 4.90366455e+00 & 7.48766737e-04 & -2.09302326e+00 \\ 7.48766737e-04 & 2.75281889e-06 & -2.18023256e-03 \\ -2.09302326e+00 & -2.18023256e-03 & 2.22674419e+00 \end{bmatrix} \begin{bmatrix} 1871 \\ 4203712 \\ 5921 \end{bmatrix} \\ &= \begin{bmatrix} -7.04346018e+01 \\ 6.38433756e-02 \\ 1.03436047e+02 \end{bmatrix} \end{aligned}$$

- ▶ Predict the value of \hat{y} , when $x_1 = 4000$ and $x_2 = 4$.

$$\hat{y} = -7.04346018e+01 + 6.38433756e-02 \times 4000 + 1.03436047e+02 \times 4 \approx 599$$



Normal Equation in TensorFlow (1/2)

```
import numpy as np
import tensorflow as tf
from sklearn.datasets import fetch_california_housing
```



Normal Equation in TensorFlow (1/2)

```
import numpy as np
import tensorflow as tf
from sklearn.datasets import fetch_california_housing

housing = fetch_california_housing()

X_train = housing.data
y_train = housing.target.reshape(-1, 1) # reshaping is done to convert y from vector to matrix
```



Normal Equation in TensorFlow (1/2)

```
import numpy as np
import tensorflow as tf
from sklearn.datasets import fetch_california_housing

housing = fetch_california_housing()

X_train = housing.data
y_train = housing.target.reshape(-1, 1) # reshaping is done to convert y from vector to matrix

# add the bias input feature i.e. a column of 1's

m = len(y_train)
X_train = np.c_[np.ones(m), X_train]
```



Normal Equation in TensorFlow (2/2)

```
# create TensorFlow Constants to store data  
  
X = tf.constant(X_train, tf.float32, name="X")  
y = tf.constant(y_train, tf.float32, name="y")
```



Normal Equation in TensorFlow (2/2)

```
# create TensorFlow Constants to store data  
  
X = tf.constant(X_train, tf.float32, name="X")  
y = tf.constant(y_train, tf.float32, name="y")
```

```
# use Normal Equation, i.e., w = (X^T.X)^-1.X.y  
  
X_T = tf.transpose(X)  
temp = tf.matrix_inverse(tf.matmul(X_T, X))  
w = tf.matmul(tf.matmul(temp, X_T), y)
```



Normal Equation in TensorFlow (2/2)

```
# create TensorFlow Constants to store data

X = tf.constant(X_train, tf.float32, name="X")
y = tf.constant(y_train, tf.float32, name="y")
```

```
# use Normal Equation, i.e., w = (X^T.X)^-1.X.y

X_T = tf.transpose(X)
temp = tf.matrix_inverse(tf.matmul(X_T, X))
w = tf.matmul(tf.matmul(temp, X_T), y)
```

```
# create TensorFlow Session

with tf.Session() as sess:
    weights = w.eval()
print(weights)
```



Normal Equation - Computational Complexity

- ▶ The computational complexity of inverting $\mathbf{X}^T \mathbf{X}$ is $O(n^3)$.
 - For an $m \times n$ matrix (where n is the number of features).

Normal Equation - Computational Complexity

- ▶ The computational complexity of inverting $\mathbf{X}^T \mathbf{X}$ is $O(n^3)$.
 - For an $m \times n$ matrix (where n is the number of features).
- ▶ But, this equation is linear with regards to the number of instances in the training set (it is $O(m)$).
 - It handles large training sets efficiently, provided they can fit in memory.



Gradient Descent





Gradient Descent (1/2)

- ▶ Gradient descent is a generic optimization algorithm capable of finding optimal solutions to a wide range of problems.



Gradient Descent (1/2)

- ▶ Gradient descent is a generic optimization algorithm capable of finding optimal solutions to a wide range of problems.
- ▶ To tweak parameters w iteratively in order to minimize a cost function $J(w)$.

Gradient Descent (2/2)

- ▶ Suppose you are **lost** in the **mountains** in a dense fog.



Gradient Descent (2/2)

- ▶ Suppose you are **lost** in the **mountains** in a dense fog.
- ▶ You can only feel the **slope** of the ground below your feet.




Gradient Descent (2/2)

- ▶ Suppose you are **lost** in the **mountains** in a dense fog.
- ▶ You can only feel the **slope** of the ground below your feet.
- ▶ A strategy to **get to the bottom** of the valley is to **go downhill** in the **direction of the steepest slope**.




Gradient Descent - Iterative Optimization Algorithm

- ▶ Choose a starting point, e.g., filling w with random values.




Gradient Descent - Iterative Optimization Algorithm

- ▶ Choose a **starting point**, e.g., filling **w** with **random values**.
- ▶ If the **stopping criterion** is true return the **current solution**, otherwise continue.




Gradient Descent - Iterative Optimization Algorithm

- ▶ Choose a **starting point**, e.g., filling w with **random values**.
- ▶ If the **stopping criterion** is true return the **current solution**, otherwise continue.
- ▶ Find a **descent direction**, a **direction in which the function value decreases** near the current point.



Gradient Descent - Iterative Optimization Algorithm

- ▶ Choose a **starting point**, e.g., filling **w** with **random values**.
- ▶ If the **stopping criterion** is true return the **current solution**, otherwise continue.
- ▶ Find a **descent direction**, a **direction in which the function value decreases** near the current point.
- ▶ Determine the **step size**, the **length of a step** in the given direction.






Gradient Descent - Key Points

- ▶ Stopping criterion
- ▶ Descent direction
- ▶ Step size (learning rate)

Gradient Descent - Stopping Criterion

- ▶ The **cost function minimum** property: the **gradient** has to be **zero**.

$$\nabla_w J(\mathbf{w}) = 0$$





Gradient Descent - Descent Direction (1/2)


- ▶ Direction in which the **function value decreases** near the current point.
- ▶ Find the **direction of descent (slope)**.

Gradient Descent - Descent Direction (1/2)

- ▶ Direction in which the **function value decreases** near the current point.
- ▶ Find the **direction of descent (slope)**.
- ▶ Example:


$$J(w) = w^2$$

$$\frac{\partial J(w)}{\partial w} = 2w = -2 \text{ at } w = -1$$



Gradient Descent - Descent Direction (2/2)

- Follow the opposite direction of the slope.






Gradient Descent - Learning Rate

- ▶ **Learning rate**: the length of steps.



Gradient Descent - Learning Rate

- ▶ **Learning rate**: the **length of steps**.
- ▶ If it is **too small**: **many iterations** to converge.



Gradient Descent - Learning Rate

- ▶ **Learning rate**: the length of steps.
- ▶ If it is **too small**: many iterations to converge.
- ▶ If it is **too high**: the algorithm might diverge.






Gradient Descent - How to Learn Model Parameters \mathbf{w} ?

- **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$.


Gradient Descent - How to Learn Model Parameters \mathbf{w} ?

- ▶ **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$.
- ▶ Start at a **random point**, and repeat the following **steps**, until the **stopping criterion** is satisfied:




Gradient Descent - How to Learn Model Parameters \mathbf{w} ?

- ▶ **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$




Gradient Descent - How to Learn Model Parameters \mathbf{w} ?

- ▶ **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$
 2. Choose a step size η




Gradient Descent - How to Learn Model Parameters \mathbf{w} ?

- ▶ Goal: find \mathbf{w} that minimizes $J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$
 2. Choose a step size η
 3. Update the parameters: $\mathbf{w}^{(\text{next})} = \mathbf{w} - \eta \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$
(should be done for all parameters simultaneously)



Gradient Descent - Different Algorithms

- ▶ Batch gradient descent
- ▶ Stochastic gradient descent
- ▶ Mini-batch gradient descent



[<https://towardsdatascience.com/gradient-descent-algorithm-and-its-variants-10f652806a3>]



Batch Gradient Descent



Batch Gradient Descent (1/2)

- ▶ Repeat the following **steps**, until the **stopping criterion** is satisfied:

1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ for all parameters \mathbf{w} .

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$



Batch Gradient Descent (1/2)

- ▶ Repeat the following **steps**, until the **stopping criterion** is satisfied:

1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ for all parameters \mathbf{w} .

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$

$$\frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{2}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}_j^{(i)}_2$$



Batch Gradient Descent (1/2)

- ▶ Repeat the following **steps**, until the **stopping criterion** is satisfied:

1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ for all parameters \mathbf{w} .

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$

$$\frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{2}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}_j^{(i)}_2$$

2. Choose a **step size** η



Batch Gradient Descent (1/2)

- ▶ Repeat the following **steps**, until the **stopping criterion** is satisfied:

1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ for all parameters \mathbf{w} .

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$

$$\frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{2}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}_j^{(i)}_2$$

2. Choose a **step size** η
3. **Update** the parameters: $w_j^{(\text{next})} = w_j - \eta \frac{\partial J(\mathbf{w})}{\partial w_j}$



Batch Gradient Descent (2/2)

- ▶ **Batch Gradient Descent:** at each step the calculation is over the **full training set \mathbf{X} .**

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$



Batch Gradient Descent (2/2)

- ▶ **Batch Gradient Descent:** at each step the calculation is over the **full training set X .**

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$

- ▶ As a result it is **slow on very large training sets**, i.e., large m .



Batch Gradient Descent (2/2)

- ▶ **Batch Gradient Descent:** at each step the calculation is over the **full training set X .**

$$J(\mathbf{w}) = \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$$

- ▶ As a result it is **slow on very large training sets**, i.e., large m .
- ▶ But, it **scales well** with the **number of features n** .

Batch Gradient Descent - Example (1/5)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| 1416 | 2 | 232 |
| 3000 | 4 | 540 |

$$\hat{y} = w_0 + w_1 x_1 + w_2 x_2$$

$$\mathbf{X} = \left[\begin{array}{ccc|c} 1 & 2104 & 3 & 400 \\ 1 & 1600 & 3 & 330 \\ 1 & 2400 & 3 & 369 \\ 1 & 1416 & 2 & 232 \\ 1 & 3000 & 4 & 540 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$

Batch Gradient Descent - Example (2/5)

$$\mathbf{X} = \left[\begin{array}{c|ccc} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$

$$\begin{aligned} \frac{\partial J(\mathbf{w})}{\partial w_0} &= \frac{2}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_0^{(i)} \\ &= \frac{2}{5} [(w_0 + 2104w_1 + 3w_2 - 400) + (w_0 + 1600w_1 + 3w_2 - 330) + \\ &\quad (w_0 + 2400w_1 + 3w_2 - 369) + (w_0 + 1416w_1 + 2w_2 - 232) + (w_0 + 3000w_1 + 4w_2 - 540)] \end{aligned}$$

Batch Gradient Descent - Example (3/5)

$$\mathbf{X} = \begin{bmatrix} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{bmatrix}$$

$$\begin{aligned}\frac{\partial J(\mathbf{w})}{\partial w_1} &= \frac{2}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}_1^{(i)} \\ &= \frac{2}{5} [2104(w_0 + 2104w_1 + 3w_2 - 400) + 1600(w_0 + 1600w_1 + 3w_2 - 330) + \\ &\quad 2400(w_0 + 2400w_1 + 3w_2 - 369) + 1416(w_0 + 1416w_1 + 2w_2 - 232) + 3000(w_0 + 3000w_1 + 4w_2 - 540)]\end{aligned}$$

Batch Gradient Descent - Example (4/5)

$$\mathbf{X} = \left[\begin{array}{c|ccc} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$

$$\begin{aligned}
 \frac{\partial J(\mathbf{w})}{\partial w_2} &= \frac{2}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_2^{(i)} \\
 &= \frac{2}{5} [3(w_0 + 2104w_1 + 3w_2 - 400) + 3(w_0 + 1600w_1 + 3w_2 - 330) + \\
 &\quad 3(w_0 + 2400w_1 + 3w_2 - 369) + 2(w_0 + 1416w_1 + 2w_2 - 232) + 4(w_0 + 3000w_1 + 4w_2 - 540)]
 \end{aligned}$$



Batch Gradient Descent - Example (5/5)

$$w_0^{(\text{next})} = w_0 - \eta \frac{\partial J(w)}{\partial w_0}$$

$$w_1^{(\text{next})} = w_1 - \eta \frac{\partial J(w)}{\partial w_1}$$

$$w_2^{(\text{next})} = w_2 - \eta \frac{\partial J(w)}{\partial w_2}$$



Stochastic Gradient Descent



Stochastic Gradient Descent

- ▶ Batch gradient descent problem: it's **slow**, because it uses the **whole training set** to compute the gradients at **every step**.



Stochastic Gradient Descent

- ▶ Batch gradient descent problem: it's **slow**, because it uses the **whole training set** to compute the gradients at **every step**.
- ▶ Stochastic gradient descent computes the gradients based on only a **single instance**.
 - It picks a **random instance** in the **training set** at **every step**.



Stochastic Gradient Descent

- ▶ Batch gradient descent problem: it's **slow**, because it uses the **whole training set** to compute the gradients at **every step**.
- ▶ Stochastic gradient descent computes the gradients based on only a **single instance**.
 - It picks a **random instance** in the **training set** at **every step**.
- ▶ The algorithm is much **faster**, but **less regular** than batch gradient descent.



Stochastic Gradient Descent - Example (1/3)

| Living area | No. of bedrooms | Price |
|-------------|-----------------|-------|
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| 1416 | 2 | 232 |
| 3000 | 4 | 540 |

$$\hat{y} = w_0 + w_1 x_1 + w_2 x_2$$

$$\mathbf{X} = \left[\begin{array}{c|ccc} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$

Stochastic Gradient Descent - Example (2/3)

$$\mathbf{X} = \begin{bmatrix} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{bmatrix}$$

$$\frac{\partial J(\mathbf{w})}{\partial w_0} = \frac{2}{m} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_0^{(i)} = \frac{2}{5} [(w_0 + 1600w_1 + 3w_2 - 330)]$$

Stochastic Gradient Descent - Example (2/3)

$$\mathbf{X} = \left[\begin{array}{c|ccc} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$

$$\frac{\partial J(\mathbf{w})}{\partial w_0} = \frac{2}{m} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_0^{(i)} = \frac{2}{5} [(w_0 + 1600w_1 + 3w_2 - 330)]$$

$$\frac{\partial J(\mathbf{w})}{\partial w_1} = \frac{2}{m} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_1^{(i)} = \frac{2}{5} [1416(w_0 + 1416w_1 + 2w_2 - 232)]$$



Stochastic Gradient Descent - Example (2/3)

$$\mathbf{X} = \left[\begin{array}{c|ccc} 1 & 2104 & 3 \\ 1 & 1600 & 3 \\ 1 & 2400 & 3 \\ 1 & 1416 & 2 \\ 1 & 3000 & 4 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 400 \\ 330 \\ 369 \\ 232 \\ 540 \end{array} \right]$$

$$\frac{\partial J(\mathbf{w})}{\partial w_0} = \frac{2}{m} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_0^{(i)} = \frac{2}{5} [(w_0 + 1600w_1 + 3w_2 - 330)]$$

$$\frac{\partial J(\mathbf{w})}{\partial w_1} = \frac{2}{m} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_1^{(i)} = \frac{2}{5} [1416(w_0 + 1416w_1 + 2w_2 - 232)]$$

$$\frac{\partial J(\mathbf{w})}{\partial w_2} = \frac{2}{m} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) x_2^{(i)} = \frac{2}{5} [3(w_0 + 2104w_1 + 3w_2 - 400)]$$

Stochastic Gradient Descent - Example (3/3)

$$w_0^{(\text{next})} = w_0 - \eta \frac{\partial J(w)}{\partial w_0}$$

$$w_1^{(\text{next})} = w_1 - \eta \frac{\partial J(w)}{\partial w_1}$$

$$w_2^{(\text{next})} = w_2 - \eta \frac{\partial J(w)}{\partial w_2}$$



Mini-Batch Gradient Descent



Mini-Batch Gradient Descent

- ▶ Batch gradient descent: at each step, it computes the gradients based on the **full training set**.



Mini-Batch Gradient Descent

- ▶ Batch gradient descent: at each step, it computes the gradients based on the **full training set**.
- ▶ Stochastic gradient descent: at each step, it computes the gradients based on **just one instance**.




Mini-Batch Gradient Descent

- ▶ **Batch gradient descent**: at each step, it computes the gradients based on the **full training set**.
- ▶ **Stochastic gradient descent**: at each step, it computes the gradients based on **just one instance**.
- ▶ **Mini-batch gradient descent**: at each step, it computes the gradients based on small **random sets of instances** called **mini-batches**.

Comparison of Algorithms for Linear Regression

| Algorithm | Large m | Large n |
|-----------------|-----------|-----------|
| Normal Equation | Fast | Slow |
| Batch GD | Slow | Fast |
| Stochastic GD | Fast | Fast |
| Mini-batch GD | Fast | Fast |





Gradient Descent in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)
```



Gradient Descent in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - y_true))
```



Gradient Descent in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - y_true))

learning_rate = 0.1
w_gradient = tf.reduce_mean((y_hat - y_true) * X) * 2
w_descent = w - learning_rate * w_gradient
w_update = tf.assign(w, w_descent)
```



Gradient Descent in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - y_true))

learning_rate = 0.1
w_gradient = tf.reduce_mean((y_hat - y_true) * X) * 2
w_descent = w - learning_rate * w_gradient
w_update = tf.assign(w, w_descent)

b_gradient = tf.reduce_mean(y_hat - y_true) * 2
b_descent = b - learning_rate * b_gradient
b_update = tf.assign(b, b_descent)
```



Gradient Descent in TensorFlow - Second Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)
```



Gradient Descent in TensorFlow - Second Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - Y))
```



Gradient Descent in TensorFlow - Second Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - Y))

learning_rate = 0.1
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
gvs = optimizer.compute_gradients(cost, [w, b])
apply_gradients = optimizer.apply_gradients(gvs)
```



Gradient Descent in TensorFlow - Third Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)
```



Gradient Descent in TensorFlow - Third Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - y_true))
```



Gradient Descent in TensorFlow - Third Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)


w = tf.Variable(5.)
b = tf.Variable(5.)

y_hat = tf.matmul(w, tf.transpose(x)) + b
cost = tf.reduce_mean(tf.square(y_hat - y_true))

learning_rate = 0.1
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
op = optimizer.minimize(cost)
```



Generalization





Training Data and Test Data

- ▶ Split data into a **training set** and a **test set**.

Full Dataset:

| | |
|---------------|-----------|
| Training Data | Test Data |
|---------------|-----------|



Training Data and Test Data

- ▶ Split data into a **training set** and a **test set**.
- ▶ Use **training set** when **training a machine learning model**.
 - Try to reduce this training error.

Full Dataset:

| | |
|---------------|-----------|
| Training Data | Test Data |
|---------------|-----------|



Training Data and Test Data

- ▶ Split data into a **training set** and a **test set**.
- ▶ Use **training set** when **training a machine learning model**.
 - Try to reduce this training error.
- ▶ Use **test set** to **measure the accuracy of the model**.
 - **Test error** is the error when you run the **trained model** on **test data (new data)**.

Full Dataset:

| | |
|---------------|-----------|
| Training Data | Test Data |
|---------------|-----------|



Generalization

- ▶ **Generalization:** make a model that performs **well** on **test data**.



Generalization

- ▶ **Generalization:** make a model that performs **well** on **test data**.
 - Have a **small test error**.



Generalization

- ▶ **Generalization:** make a model that performs **well** on **test data**.
 - Have a **small test error**.
- ▶ **Challenges**
 1. Make the **training error small**.
 2. Make the **gap** between **training** and **test error small**.



More About The Test Error

- ▶ The **test error** is computed as the **MSE** of **k** test instances.

$$\text{MSE}_{\text{test}} = \frac{1}{k} \sum_i^k (\hat{y}_{\text{test}}^{(i)} - y_{\text{test}}^{(i)})^2 = E[(\hat{y}_{\text{test}} - y_{\text{test}})^2]$$



More About The Test Error

- ▶ The **test error** is computed as the **MSE** of **k** test instances.

$$\text{MSE}_{\text{test}} = \frac{1}{k} \sum_i^k (\hat{y}_{\text{test}}^{(i)} - y_{\text{test}}^{(i)})^2 = E[(\hat{y}_{\text{test}} - y_{\text{test}})^2]$$

- ▶ A model's **test error** can be expressed as the **sum** of **bias** and **variance**.

$$E[(\hat{y}_{\text{test}} - y_{\text{test}})^2] = \text{Bias}[\hat{y}_{\text{test}}, y_{\text{test}}]^2 + \text{Var}[\hat{y}_{\text{test}}] + \varepsilon^2$$



Bias and Underfitting

- ▶ Bias: the expected deviation from the true value of the function.


$$\text{Bias}[\hat{y}_{\text{test}}, y_{\text{test}}] = E[\hat{y}_{\text{test}}] - y_{\text{test}}$$

Bias and Underfitting

- ▶ Bias: the expected deviation from the true value of the function.

$$\text{Bias}[\hat{y}_{\text{test}}, y_{\text{test}}] = E[\hat{y}_{\text{test}}] - y_{\text{test}}$$

- ▶ A high-bias model is most likely to underfit the training data.
 - High error value on the training set.




Bias and Underfitting

- ▶ Bias: the expected deviation from the true value of the function.

$$\text{Bias}[\hat{y}_{\text{test}}, y_{\text{test}}] = E[\hat{y}_{\text{test}}] - y_{\text{test}}$$

- ▶ A high-bias model is most likely to underfit the training data.
 - High error value on the training set.
- ▶ Underfitting happens when the model is too simple to learn the underlying structure of the data.





Variance and Overfitting

- ▶ **Variance**: how much a model changes if you train it on a different training set.


$$\text{Var}[\hat{y}_{\text{test}}] = E[(\hat{y}_{\text{test}} - E[\hat{y}_{\text{test}}])^2]$$

Variance and Overfitting

- ▶ **Variance**: how much a model changes if you train it on a different training set.

$$\text{Var}[\hat{y}_{\text{test}}] = E[(\hat{y}_{\text{test}} - E[\hat{y}_{\text{test}}])^2]$$

- ▶ A **high-variance** model is most likely to **overfit** the training data.
 - The **gap** between the **training error** and **test error** is **too large**.




Variance and Overfitting


- ▶ **Variance**: how much a model changes if you train it on a different training set.

$$\text{Var}[\hat{y}_{\text{test}}] = E[(\hat{y}_{\text{test}} - E[\hat{y}_{\text{test}}])^2]$$

- ▶ A **high-variance** model is most likely to **overfit** the training data.
 - The **gap** between the **training error** and **test error** is **too large**.
- ▶ **Overfitting** happens when the **model is too complex** relative to the amount and noisiness of the training data.



Underfitting



Optimal



Overfitting



The Bias/Variance Tradeoff (1/2)

- ▶ Assume a model with two parameters w_0 (intercept) and w_1 (slope): $\hat{y} = w_0 + w_1 x$



The Bias/Variance Tradeoff (1/2)

- ▶ Assume a model with **two parameters** w_0 (**intercept**) and w_1 (**slope**): $\hat{y} = w_0 + w_1 x$
- ▶ They give the learning algorithm **two degrees of freedom**.



The Bias/Variance Tradeoff (1/2)

- ▶ Assume a model with two parameters w_0 (intercept) and w_1 (slope): $\hat{y} = w_0 + w_1x$
- ▶ They give the learning algorithm two degrees of freedom.
- ▶ We tweak both the w_0 and w_1 to adapt the model to the training data.




The Bias/Variance Tradeoff (1/2)

- ▶ Assume a model with two parameters w_0 (intercept) and w_1 (slope): $\hat{y} = w_0 + w_1 x$
- ▶ They give the learning algorithm two degrees of freedom.
- ▶ We tweak both the w_0 and w_1 to adapt the model to the training data.
- ▶ If we forced $w_0 = 0$, the algorithm would have only one degree of freedom and would have a much harder time fitting the data properly.

The Bias/Variance Tradeoff (2/2)

- ▶ Increasing degrees of freedom will typically increase its variance and reduce its bias.
- ▶ Decreasing degrees of freedom increases its bias and reduces its variance.
- ▶ This is why it is called a **tradeoff**.



© Machine Learning @ Berkeley

[<https://ml.berkeley.edu/blog/2017/07/13/tutorial-4>]



Regularization (1/2)

- ▶ One way to reduce the risk of overfitting is to have fewer degrees of freedom.
- ▶ Regularization is a technique to reduce the risk of overfitting.
- ▶ For a linear model, regularization is achieved by constraining the weights of the model.

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \lambda R(\mathbf{w})$$



Regularization (2/2)

- ▶ Lasso regression (1): $R(\mathbf{w}) = \lambda \sum_{i=1}^n |w_i|$ is added to the cost function:

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \lambda \sum_{i=1}^n |w_i|$$



Regularization (2/2)

- ▶ Lasso regression (1): $R(\mathbf{w}) = \lambda \sum_{i=1}^n |w_i|$ is added to the cost function:

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \lambda \sum_{i=1}^n |w_i|$$

- ▶ Ridge regression (2): $R(\mathbf{w}) = \lambda \sum_{i=1}^n w_i^2$ is added to the cost function.

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \lambda \sum_{i=1}^n w_i^2$$



Regularization (2/2)

- ▶ Lasso regression (/1): $R(\mathbf{w}) = \lambda \sum_{i=1}^n |w_i|$ is added to the cost function:

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \lambda \sum_{i=1}^n |w_i|$$

- ▶ Ridge regression (/2): $R(\mathbf{w}) = \lambda \sum_{i=1}^n w_i^2$ is added to the cost function.

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \lambda \sum_{i=1}^n w_i^2$$

- ▶ ElasticNet: a middle ground between /1 and /2 regularization.

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) + \alpha \lambda \sum_{i=1}^n |w_i| + (1 - \alpha) \lambda \sum_{i=1}^n w_i^2$$



Hyperparameters



Hyperparameters and Validation Sets (1/2)

- ▶ Hyperparameters are settings that we can use to control the behavior of a learning algorithm.



Hyperparameters and Validation Sets (1/2)

- ▶ Hyperparameters are settings that we can use to control the behavior of a learning algorithm.
- ▶ The values of hyperparameters are not adapted by the learning algorithm itself.
 - E.g., the α and λ values for regularization.



Hyperparameters and Validation Sets (1/2)

- ▶ Hyperparameters are settings that we can use to control the behavior of a learning algorithm.
- ▶ The values of hyperparameters are not adapted by the learning algorithm itself.
 - E.g., the α and λ values for regularization.
- ▶ We do not learn the hyperparameter.
 - It is not appropriate to learn that hyperparameter on the training set.
 - If learned on the training set, such hyperparameters would always result in overfitting.



Hyperparameters and Validation Sets (2/2)

- ▶ To find **hyperparameters**, we need a **validation set** of examples that the **training algorithm does not observe**.



Hyperparameters and Validation Sets (2/2)

- ▶ To find **hyperparameters**, we need a **validation set** of examples that the **training algorithm does not observe**.
- ▶ We construct the **validation set** from the **training data** (**not the test data**).



Hyperparameters and Validation Sets (2/2)

- ▶ To find **hyperparameters**, we need a **validation set** of examples that the **training algorithm does not observe**.
- ▶ We construct the **validation set** from the **training data** (**not the test data**).
- ▶ We split the **training data** into two disjoint subsets:
 1. One is used to **learn the parameters**.
 2. The other one (the **validation set**) is used to **estimate the test error** **during or after training**, allowing for the **hyperparameters** to be updated accordingly.

Full Dataset:

| Training Data | Validation Data | Test Data |
|---------------|-----------------|-----------|
| | | |


Cross-Validation

- ▶ **Cross-validation:** a technique to avoid **wasting too much training data in validation sets.**



Cross-Validation

- ▶ **Cross-validation:** a technique to avoid **wasting too much training data in validation sets.**
- ▶ The **training set** is split into **complementary subsets**.



Cross-Validation

- ▶ **Cross-validation:** a technique to avoid **wasting too much training data** in **validation sets**.
- ▶ The **training set** is split into **complementary subsets**.
- ▶ Each model is **trained** against a different **combination** of these subsets and **validated** against the **remaining parts**.





Logistic Regression



Let's Start with an Example

Example (1/4)

- ▶ Given the dataset of m cancer tests.

| Tumor size | Cancer |
|------------|--------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| : | : |

Example (1/4)

- Given the dataset of m cancer tests.


| Tumor size | Cancer |
|------------|--------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| : | : |

- Predict the risk of cancer, as a function of the tumor size?

Example (2/4)

| Tumor size | Cancer |
|------------|--------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| : | : |
| : | : |


$$\mathbf{x} = \begin{bmatrix} 330 \\ 120 \\ 400 \\ \vdots \\ \vdots \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \\ \vdots \end{bmatrix}$$



Example (2/4)

| Tumor size | Cancer |
|------------|--------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| : | : |
| : | : |

$$\mathbf{x} = \begin{bmatrix} 330 \\ 120 \\ 400 \\ \vdots \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \end{bmatrix}$$




- $\mathbf{x}^{(i)} \in \mathbb{R}$: $x_1^{(i)}$ is the tumor size of the i th instance in the training set.

Example (3/4)

| Tumor size | Cancer |
|------------|--------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| ⋮ | ⋮ |

$$x = \begin{bmatrix} 330 \\ 120 \\ 400 \\ \vdots \end{bmatrix} \quad y = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \end{bmatrix}$$




- ▶ Predict the risk of cancer \hat{y} as a function of the tumor sizes x_1 , i.e., $\hat{y} = f(x_1)$
- ▶ E.g., what is \hat{y} , if $x_1 = 500$?

Example (3/4)

| Tumor size | Cancer |
|------------|----------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| \vdots | \vdots |

$$x = \begin{bmatrix} 330 \\ 120 \\ 400 \\ \vdots \end{bmatrix} \quad y = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \end{bmatrix}$$




- Predict the risk of cancer \hat{y} as a function of the tumor sizes x_1 , i.e., $\hat{y} = f(x_1)$
- E.g., what is \hat{y} , if $x_1 = 500$?
- As an initial choice: $\hat{y} = f_w(x) = w_0 + w_1 x_1$

Example (3/4)



| Tumor size | Cancer |
|------------|----------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |
| \vdots | \vdots |

$$x = \begin{bmatrix} 330 \\ 120 \\ 400 \\ \vdots \end{bmatrix} \quad y = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \end{bmatrix}$$



- ▶ Predict the risk of cancer \hat{y} as a function of the tumor sizes x_1 , i.e., $\hat{y} = f(x_1)$
- ▶ E.g., what is \hat{y} , if $x_1 = 500$?
- ▶ As an initial choice: $\hat{y} = f_w(x) = w_0 + w_1 x_1$
- ▶ Bad model!

Example (4/4)




- ▶ A better model $\hat{y} = \frac{1}{1+e^{-(w_0+w_1x_1)}}$

Sigmoid Function

- The **sigmoid function**, denoted by $\sigma(\cdot)$, outputs a number **between 0 and 1**.

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$




- When $t < 0$, then $\sigma(t) < 0.5$
- when $t \geq 0$, then $\sigma(t) \geq 0.5$



Binomial Logistic Regression

Binomial Logistic Regression (1/2)

- ▶ Our goal: to build a system that takes input $\mathbf{x} \in \mathbb{R}^n$ and predicts output $\hat{\mathbf{y}} \in \{0, 1\}$.
- ▶ To specify which of 2 categories an input \mathbf{x} belongs to.





Binomial Logistic Regression (2/2)

- ▶ Linear regression

$$\hat{y} = w_0x_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n = \mathbf{w}^T\mathbf{x}$$

Binomial Logistic Regression (2/2)

- ▶ Linear regression

$$\hat{y} = w_0x_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n = \mathbf{w}^T\mathbf{x}$$

- ▶ Binomial logistic regression


$$z = w_0x_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n = \mathbf{w}^T\mathbf{x}$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-\mathbf{w}^T\mathbf{x}}}$$



How to Learn Model Parameters w ?

Linear Regression - Cost Function



- One reasonable model should make \hat{y} close to y , at least for the training dataset.
- Cost function $J(\mathbf{w})$: the mean squared error (MSE)

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = (\hat{y}^{(i)} - y^{(i)})^2$$

$$J(\mathbf{w}) = \frac{1}{m} \sum_i^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_i^m (\hat{y}^{(i)} - y^{(i)})^2$$



Binomial Logistic Regression - Cost Function (1/5)

- ▶ Naive idea: minimizing the Mean Squared Error (MSE)

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = (\hat{y}^{(i)} - y^{(i)})^2$$

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$



Binomial Logistic Regression - Cost Function (1/5)

- ▶ Naive idea: minimizing the Mean Squared Error (MSE)

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = (\hat{y}^{(i)} - y^{(i)})^2$$

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \left(\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}^{(i)}}} - y^{(i)} \right)^2$$

Binomial Logistic Regression - Cost Function (1/5)

- ▶ Naive idea: minimizing the Mean Squared Error (MSE)

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = (\hat{y}^{(i)} - y^{(i)})^2$$


$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

$$J(\mathbf{w}) = \text{MSE}(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \left(\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}^{(i)}}} - y^{(i)} \right)^2$$

- ▶ This cost function is a non-convex function for parameter optimization.


Binomial Logistic Regression - Cost Function (2/5)

- ▶ What do we mean by **non-convex**?
- ▶ If a line joining two points on the curve, **crosses the curve**.
- ▶ The algorithm may converge to a **local minimum**.



Binomial Logistic Regression - Cost Function (2/5)

- ▶ What do we mean by **non-convex**?
- ▶ If a line joining two points on the curve, **crosses the curve**.
- ▶ The algorithm may converge to a **local minimum**.
- ▶ We want a **convex** logistic regression **cost function $J(\mathbf{w})$** .





Binomial Logistic Regression - Cost Function (3/5)

- ▶ The predicted value $\hat{y} = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$
- ▶ $\text{cost}(\hat{y}^{(i)}, y^{(i)}) = ?$



Binomial Logistic Regression - Cost Function (3/5)

- ▶ The predicted value $\hat{y} = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$
- ▶ $\text{cost}(\hat{y}^{(i)}, y^{(i)}) = ?$
- ▶ The $\text{cost}(\hat{y}^{(i)}, y^{(i)})$ should be
 - Close to 0, if the predicted value \hat{y} will be close to true value y .
 - Large, if the predicted value \hat{y} will be far from the true value y .



Binomial Logistic Regression - Cost Function (3/5)

- ▶ The predicted value $\hat{y} = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$
- ▶ $\text{cost}(\hat{y}^{(i)}, y^{(i)}) = ?$
- ▶ The $\text{cost}(\hat{y}^{(i)}, y^{(i)})$ should be
 - Close to 0, if the predicted value \hat{y} will be close to true value y .
 - Large, if the predicted value \hat{y} will be far from the true value y .

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = \begin{cases} -\log(\hat{y}^{(i)}) & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

Binomial Logistic Regression - Cost Function (4/5)

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = \begin{cases} -\log(\hat{y}^{(i)}) & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$





Binomial Logistic Regression - Cost Function (5/5)

- We can define $J(\mathbf{w})$ as below

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = \begin{cases} -\log(\hat{y}^{(i)}) & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

Binomial Logistic Regression - Cost Function (5/5)

- We can define $J(\mathbf{w})$ as below

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = \begin{cases} -\log(\hat{y}^{(i)}) & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$



How to Learn Model Parameters w ?

- ▶ We want to choose w so as to minimize $J(w)$.
- ▶ An approach to find w : gradient descent
 - Batch gradient descent
 - Stochastic gradient descent
 - Mini-batch gradient descent



Binomial Logistic Regression Gradient Descent (1/2)

- ▶ Goal: find \mathbf{w} that minimizes $J(\mathbf{w}) = -\frac{1}{m} \sum_i^m (y^{(i)} \log(\hat{y}^{(i)}) + (1-y^{(i)}) \log(1-\hat{y}^{(i)}))$.



Binomial Logistic Regression Gradient Descent (1/2)

- ▶ **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = -\frac{1}{m} \sum_i^m (y^{(i)} \log(\hat{y}^{(i)}) + (1-y^{(i)}) \log(1-\hat{y}^{(i)}))$.
- ▶ Start at a **random point**, and repeat the following **steps**, until the **stopping criterion** is satisfied:



Binomial Logistic Regression Gradient Descent (1/2)

- ▶ **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = -\frac{1}{m} \sum_i^m (y^{(i)} \log(\hat{y}^{(i)}) + (1-y^{(i)}) \log(1-\hat{y}^{(i)}))$.
- ▶ Start at a **random point**, and repeat the following **steps**, until the **stopping criterion** is satisfied:
 1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$



Binomial Logistic Regression Gradient Descent (1/2)

- ▶ **Goal:** find \mathbf{w} that minimizes $J(\mathbf{w}) = -\frac{1}{m} \sum_i^m (y^{(i)} \log(\hat{y}^{(i)}) + (1-y^{(i)}) \log(1-\hat{y}^{(i)}))$.
- ▶ Start at a **random point**, and repeat the following **steps**, until the **stopping criterion** is satisfied:
 1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$
 2. Choose a **step size** η



Binomial Logistic Regression Gradient Descent (1/2)

- ▶ Goal: find \mathbf{w} that minimizes $J(\mathbf{w}) = -\frac{1}{m} \sum_i^m (y^{(i)} \log(\hat{y}^{(i)}) + (1-y^{(i)}) \log(1-\hat{y}^{(i)}))$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$
 2. Choose a step size η
 3. Update the parameters: $\mathbf{w}^{(\text{next})} = \mathbf{w} - \eta \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ (simultaneously for all parameters)

Binomial Logistic Regression Gradient Descent (2/2)

- ▶ 1. Determine a descent direction $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$.

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

$$\frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j$$

Binomial Logistic Regression Gradient Descent (2/2)

- ▶ 1. Determine a descent direction $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$.

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

$$\frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j$$

- ▶ 2. Choose a step size η

Binomial Logistic Regression Gradient Descent (2/2)

- ▶ 1. Determine a **descent direction** $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$.

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

$$\frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j$$

- ▶ 2. Choose a **step size** η
- ▶ 3. Update the parameters: $w_j^{(\text{next})} = w_j - \eta \frac{\partial J(\mathbf{w})}{\partial w_j}$
 - $0 \leq j \leq n$, where **n** is the **number of features**.

Binomial Logistic Regression Gradient Descent - Example (1/4)

| Tumor size | Cancer |
|------------|--------|
| 330 | 1 |
| 120 | 0 |
| 400 | 1 |

$$\mathbf{X} = \begin{bmatrix} 1 & 330 \\ 1 & 120 \\ 1 & 400 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$

- ▶ Predict the risk of cancer \hat{y} as a function of the tumor sizes x_1 .
- ▶ E.g., what is \hat{y} , if $x_1 = 500$?

Binomial Logistic Regression Gradient Descent - Example (2/4)

$$\mathbf{X} = \left[\begin{array}{c|c} 1 & 330 \\ 1 & 120 \\ 1 & 400 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 1 \\ 0 \\ 1 \end{array} \right]$$

$$\hat{y} = \sigma(w_0 + w_1 x_1) = \frac{1}{1 + e^{-(w_0 + w_1 x_1)}}$$

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

Binomial Logistic Regression Gradient Descent - Example (2/4)

$$\mathbf{X} = \left[\begin{array}{c|c} 1 & 330 \\ 1 & 120 \\ 1 & 400 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 1 \\ 0 \\ 1 \end{array} \right]$$

$$\hat{y} = \sigma(w_0 + w_1 x_1) = \frac{1}{1 + e^{-(w_0 + w_1 x_1)}}$$

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

$$\begin{aligned} \frac{\partial J(\mathbf{w})}{\partial w_0} &= \frac{1}{3} \sum_{i=1}^3 (\hat{y}^{(i)} - y^{(i)}) x_0 \\ &= \frac{1}{3} \left[\left(\frac{1}{1 + e^{-(w_0 + 330w_1)}} - 1 \right) + \left(\frac{1}{1 + e^{-(w_0 + 120w_1)}} - 0 \right) + \left(\frac{1}{1 + e^{-(w_0 + 400w_1)}} - 1 \right) \right] \end{aligned}$$

Binomial Logistic Regression Gradient Descent - Example (3/4)

$$\mathbf{X} = \left[\begin{array}{c|c} 1 & 330 \\ 1 & 120 \\ 1 & 400 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 1 \\ 0 \\ 1 \end{array} \right]$$

$$\hat{y} = \sigma(w_0 + w_1 x_1) = \frac{1}{1 + e^{-(w_0 + w_1 x_1)}}$$

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

Binomial Logistic Regression Gradient Descent - Example (3/4)

$$\mathbf{X} = \left[\begin{array}{c|c} 1 & 330 \\ 1 & 120 \\ 1 & 400 \end{array} \right] \quad \mathbf{y} = \left[\begin{array}{c} 1 \\ 0 \\ 1 \end{array} \right]$$

$$\hat{y} = \sigma(w_0 + w_1 x_1) = \frac{1}{1 + e^{-(w_0 + w_1 x_1)}}$$

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

$$\begin{aligned} \frac{\partial J(\mathbf{w})}{\partial w_1} &= \frac{1}{3} \sum_{i=1}^3 (\hat{y}^{(i)} - y^{(i)}) x_1 \\ &= \frac{1}{3} [330 \left(\frac{1}{1 + e^{-(w_0 + 330w_1)}} - 1 \right) + 120 \left(\frac{1}{1 + e^{-(w_0 + 120w_1)}} - 0 \right) + 400 \left(\frac{1}{1 + e^{-(w_0 + 400w_1)}} - 1 \right)] \end{aligned}$$



Binomial Logistic Regression Gradient Descent - Example (4/4)

$$w_0^{(\text{next})} = w_0 - \eta \frac{\partial J(\mathbf{w})}{\partial w_0}$$

$$w_1^{(\text{next})} = w_1 - \eta \frac{\partial J(\mathbf{w})}{\partial w_1}$$



Logistic Regression in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)
```



Logistic Regression in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

z = tf.matmul(w, tf.transpose(x)) + b
y_hat = tf.sigmoid(z)

cost = -y_true * tf.log(y_hat) - (1 - y_true) * tf.log(1 - y_hat)
cost = tf.reduce_mean(cost)
```



Logistic Regression in TensorFlow - First Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

z = tf.matmul(w, tf.transpose(x)) + b
y_hat = tf.sigmoid(z)

cost = -y_true * tf.log(y_hat) - (1 - y_true) * tf.log(1 - y_hat)
cost = tf.reduce_mean(cost)

learning_rate = 0.1
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
op = optimizer.minimize(cost)
```



Logistic Regression in TensorFlow - Second Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)
```



Logistic Regression in TensorFlow - Second Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

z = tf.matmul(w, tf.transpose(x)) + b
y_hat = tf.sigmoid(z)

cost = tf.nn.sigmoid_cross_entropy_with_logits(labels=y_true, logits=z)
cost = tf.reduce_mean(cost)
```



Logistic Regression in TensorFlow - Second Implementation

```
x_train = [1, 2, 3]
y_train = [1, 2, 3]

X = tf.placeholder(tf.float32)
y_true = tf.placeholder(tf.float32)

w = tf.Variable(5.)
b = tf.Variable(5.)

z = tf.matmul(w, tf.transpose(x)) + b
y_hat = tf.sigmoid(z)

cost = tf.nn.sigmoid_cross_entropy_with_logits(labels=y_true, logits=z)
cost = tf.reduce_mean(cost)

learning_rate = 0.1
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
op = optimizer.minimize(cost)
```



Multinomial Logistic Regression



Multinomial Logistic Regression

- ▶ Multinomial classifiers can distinguish between more than two classes.
- ▶ Instead of $y \in \{0, 1\}$, we have $y \in \{1, 2, \dots, k\}$.



Binomial vs. Multinomial Logistic Regression (1/2)

- ▶ In a **binomial classifier**, $y \in \{0, 1\}$, the **estimator** is $\hat{y} = p(y = 1 | x; w)$.
 - We find **one** set of parameters **w**.

$$w^T = [w_0, w_1, \dots, w_n]$$



Binomial vs. Multinomial Logistic Regression (1/2)

- ▶ In a **binomial classifier**, $y \in \{0, 1\}$, the **estimator** is $\hat{y} = p(y = 1 | \mathbf{x}; \mathbf{w})$.
 - We find **one** set of parameters \mathbf{w} .

$$\mathbf{w}^T = [w_0, w_1, \dots, w_n]$$

- ▶ In **multinomial classifier**, $y \in \{1, 2, \dots, k\}$, we need to estimate the result for each **individual label**, i.e., $\hat{y}_j = p(y = j | \mathbf{x}; \mathbf{w})$.

Binomial vs. Multinomial Logistic Regression (1/2)

- In a **binomial classifier**, $y \in \{0, 1\}$, the **estimator** is $\hat{y} = p(y = 1 | \mathbf{x}; \mathbf{w})$.
 - We find **one** set of parameters \mathbf{w} .

$$\mathbf{w}^T = [w_0, w_1, \dots, w_n]$$

- In **multinomial classifier**, $y \in \{1, 2, \dots, k\}$, we need to estimate the result for each **individual label**, i.e., $\hat{y}_j = p(y = j | \mathbf{x}; \mathbf{w})$.
 - We find **k** set of parameters \mathbf{W} .

$$\mathbf{W} = \begin{bmatrix} [w_{0,1}, w_{1,1}, \dots, w_{n,1}] \\ [w_{0,2}, w_{1,2}, \dots, w_{n,2}] \\ \vdots \\ [w_{0,k}, w_{1,k}, \dots, w_{n,k}] \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_k^T \end{bmatrix}$$



Binomial vs. Multinomial Logistic Regression (2/2)

- ▶ In a **binary class**, $y \in \{0, 1\}$, we use the **sigmoid** function.

$$\mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + \cdots + w_n x_n$$

$$\hat{y} = p(y = 1 \mid \mathbf{x}; \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

Binomial vs. Multinomial Logistic Regression (2/2)

- ▶ In a **binary class**, $y \in \{0, 1\}$, we use the **sigmoid** function.

$$\mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + \cdots + w_n x_n$$

$$\hat{y} = p(y = 1 \mid \mathbf{x}; \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

- ▶ In **multiclasses**, $y \in \{1, 2, \dots, k\}$, we use the **softmax** function.

$$\mathbf{w}_j^T \mathbf{x} = w_{0,j} x_0 + w_{1,j} x_1 + \cdots + w_{n,j} x_n, 1 \leq j \leq k$$

$$\hat{y}_j = p(y = j \mid \mathbf{x}; \mathbf{w}_j) = \sigma(\mathbf{w}_j^T \mathbf{x}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}}}{\sum_{i=1}^k e^{\mathbf{w}_i^T \mathbf{x}}}$$

Sigmoid vs. Softmax

- Sigmoid function: $\sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$



Sigmoid vs. Softmax

- ▶ Sigmoid function: $\sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$
- ▶ Softmax function: $\sigma(\mathbf{w}_j^T \mathbf{x}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}}}{\sum_{i=1}^k e^{\mathbf{w}_i^T \mathbf{x}}}$



Sigmoid vs. Softmax

- ▶ Sigmoid function: $\sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$
- ▶ Softmax function: $\sigma(\mathbf{w}_j^T \mathbf{x}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}}}{\sum_{i=1}^k e^{\mathbf{w}_i^T \mathbf{x}}}$
 - Calculate the probabilities of each target class over all possible target classes.



Sigmoid vs. Softmax

- ▶ Sigmoid function: $\sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$
- ▶ Softmax function: $\sigma(\mathbf{w}_j^T \mathbf{x}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}}}{\sum_{i=1}^k e^{\mathbf{w}_i^T \mathbf{x}}}$
 - Calculate the probabilities of each target class over all possible target classes.
 - The softmax function for two classes is equivalent the sigmoid function.



Softmax Model Estimation and Prediction - Example (1/2)

- ▶ Assume we have a **training set** consisting of $m = 4$ instances from $k = 3$ **classes**.

$$\mathbf{x}^{(1)} \rightarrow \text{class1}, \mathbf{y}^{(1)\top} = [1 \ 0 \ 0]$$

$$\mathbf{x}^{(2)} \rightarrow \text{class2}, \mathbf{y}^{(2)\top} = [0 \ 1 \ 0]$$

$$\mathbf{x}^{(3)} \rightarrow \text{class3}, \mathbf{y}^{(3)\top} = [0 \ 0 \ 1]$$

$$\mathbf{x}^{(4)} \rightarrow \text{class3}, \mathbf{y}^{(4)\top} = [0 \ 0 \ 1]$$

$$\mathbf{Y} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

Softmax Model Estimation and Prediction - Example (1/2)

- ▶ Assume we have a **training set** consisting of $m = 4$ instances from $k = 3$ **classes**.

$$\mathbf{x}^{(1)} \rightarrow \text{class1}, \mathbf{y}^{(1)\top} = [1 \ 0 \ 0]$$

$$\mathbf{x}^{(2)} \rightarrow \text{class2}, \mathbf{y}^{(2)\top} = [0 \ 1 \ 0]$$

$$\mathbf{x}^{(3)} \rightarrow \text{class3}, \mathbf{y}^{(3)\top} = [0 \ 0 \ 1]$$

$$\mathbf{x}^{(4)} \rightarrow \text{class3}, \mathbf{y}^{(4)\top} = [0 \ 0 \ 1]$$

$$\mathbf{Y} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

- ▶ Assume **training set** \mathbf{X} and random parameters \mathbf{W} are as below:

$$\mathbf{X} = \left[\begin{array}{c|ccc} 1 & 0.1 & 0.5 \\ 1 & 1.1 & 2.3 \\ 1 & -1.1 & -2.3 \\ 1 & -1.5 & -2.5 \end{array} \right] \quad \mathbf{W} = \begin{bmatrix} 0.01 & 0.1 & 0.1 \\ 0.1 & 0.2 & 0.3 \\ 0.1 & 0.2 & 0.3 \end{bmatrix}$$



Softmax Model Estimation and Prediction - Example (2/2)

- ▶ Now, let's compute the softmax activation:

$$\hat{y}_j^{(i)} = p(y^{(i)} = j \mid \mathbf{x}^{(i)}; \mathbf{w}_j) = \sigma(\mathbf{w}_j^T \mathbf{x}^{(i)}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}^{(i)}}}{\sum_{l=1}^k e^{\mathbf{w}_l^T \mathbf{x}^{(i)}}}$$

Softmax Model Estimation and Prediction - Example (2/2)

- Now, let's compute the softmax activation:

$$\hat{y}_j^{(i)} = p(y^{(i)} = j \mid \mathbf{x}^{(i)}; \mathbf{w}_j) = \sigma(\mathbf{w}_j^T \mathbf{x}^{(i)}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}^{(i)}}}{\sum_{l=1}^k e^{\mathbf{w}_l^T \mathbf{x}^{(i)}}}$$

$$\hat{\mathbf{Y}} = \begin{bmatrix} \hat{y}^{(1)\top} \\ \hat{y}^{(2)\top} \\ \hat{y}^{(3)\top} \\ \hat{y}^{(4)\top} \end{bmatrix} = \begin{bmatrix} 0.29 & 0.34 & 0.36 \\ 0.21 & 0.33 & 0.46 \\ 0.43 & 0.33 & 0.24 \\ 0.45 & 0.33 & 0.22 \end{bmatrix} \quad \text{the predicted classes} = \begin{bmatrix} 3 \\ 3 \\ 1 \\ 1 \end{bmatrix} \quad \text{The correct classes} = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 3 \end{bmatrix}$$

- They are terribly wrong.

Softmax Model Estimation and Prediction - Example (2/2)

- Now, let's compute the softmax activation:

$$\hat{y}_j^{(i)} = p(y^{(i)} = j \mid \mathbf{x}^{(i)}; \mathbf{w}_j) = \sigma(\mathbf{w}_j^T \mathbf{x}^{(i)}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}^{(i)}}}{\sum_{l=1}^k e^{\mathbf{w}_l^T \mathbf{x}^{(i)}}}$$

$$\hat{\mathbf{Y}} = \begin{bmatrix} \hat{y}^{(1)\top} \\ \hat{y}^{(2)\top} \\ \hat{y}^{(3)\top} \\ \hat{y}^{(4)\top} \end{bmatrix} = \begin{bmatrix} 0.29 & 0.34 & 0.36 \\ 0.21 & 0.33 & 0.46 \\ 0.43 & 0.33 & 0.24 \\ 0.45 & 0.33 & 0.22 \end{bmatrix} \quad \text{the predicted classes} = \begin{bmatrix} 3 \\ 3 \\ 1 \\ 1 \end{bmatrix} \quad \text{The correct classes} = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 3 \end{bmatrix}$$

- They are **terribly wrong**.
- We need to **update the weights** based on the cost function.

Softmax Model Estimation and Prediction - Example (2/2)

- Now, let's compute the softmax activation:

$$\hat{y}_j^{(i)} = p(y^{(i)} = j \mid \mathbf{x}^{(i)}; \mathbf{w}_j) = \sigma(\mathbf{w}_j^T \mathbf{x}^{(i)}) = \frac{e^{\mathbf{w}_j^T \mathbf{x}^{(i)}}}{\sum_{l=1}^k e^{\mathbf{w}_l^T \mathbf{x}^{(i)}}}$$

$$\hat{\mathbf{Y}} = \begin{bmatrix} \hat{y}^{(1)\top} \\ \hat{y}^{(2)\top} \\ \hat{y}^{(3)\top} \\ \hat{y}^{(4)\top} \end{bmatrix} = \begin{bmatrix} 0.29 & 0.34 & 0.36 \\ 0.21 & 0.33 & 0.46 \\ 0.43 & 0.33 & 0.24 \\ 0.45 & 0.33 & 0.22 \end{bmatrix} \quad \text{the predicted classes} = \begin{bmatrix} 3 \\ 3 \\ 1 \\ 1 \end{bmatrix} \quad \text{The correct classes} = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 3 \end{bmatrix}$$

- They are **terribly wrong**.
- We need to **update the weights** based on the cost function.
- What is the cost function?**



Multinomial Logistic Regression - Cost Function (1/2)

- ▶ The **objective** is to have a model that estimates a **high probability** for the target class, and consequently a **low probability** for the other classes.

Multinomial Logistic Regression - Cost Function (1/2)

- ▶ The **objective** is to have a model that estimates a **high probability** for the target class, and consequently a **low probability** for the other classes.
- ▶ **Cost function:** the **cross-entropy** between the **correct classes** and **predicted class** for all classes.

$$J(\mathbf{w}_j) = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^k y_j^{(i)} \log(\hat{y}_j^{(i)})$$

Multinomial Logistic Regression - Cost Function (1/2)

- ▶ The **objective** is to have a model that estimates a **high probability** for the target class, and consequently a **low probability** for the other classes.
- ▶ **Cost function:** the **cross-entropy** between the **correct classes** and **predicted class** for all classes.

$$J(\mathbf{w}_j) = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^k y_j^{(i)} \log(\hat{y}_j^{(i)})$$

- ▶ $y_j^{(i)}$ is 1 if the target class for the i th instance is j , otherwise, it is 0.

Multinomial Logistic Regression - Cost Function (2/2)

- ▶ If there are two classes ($k = 2$), this cost function is equivalent to the logistic regression's cost function.

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$



How to Learn Model Parameters \mathbf{W} ?

- ▶ Goal: find \mathbf{W} that minimizes $J(\mathbf{W})$.



How to Learn Model Parameters \mathbf{W} ?

- ▶ Goal: find \mathbf{W} that minimizes $J(\mathbf{W})$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:



How to Learn Model Parameters \mathbf{W} ?

- ▶ Goal: find \mathbf{W} that minimizes $J(\mathbf{W})$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{W})}{\partial \mathbf{w}}$



How to Learn Model Parameters \mathbf{W} ?

- ▶ Goal: find \mathbf{W} that minimizes $J(\mathbf{W})$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{W})}{\partial \mathbf{w}}$
 2. Choose a step size η



How to Learn Model Parameters \mathbf{W} ?

- ▶ Goal: find \mathbf{W} that minimizes $J(\mathbf{W})$.
- ▶ Start at a random point, and repeat the following steps, until the stopping criterion is satisfied:
 1. Determine a descent direction $\frac{\partial J(\mathbf{W})}{\partial \mathbf{w}}$
 2. Choose a step size η
 3. Update the parameters: $\mathbf{w}^{(\text{next})} = \mathbf{w} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{w}}$ (simultaneously for all parameters)



Performance Measures



Evaluation of Classification Models (1/3)

- ▶ In a **classification problem**, there exists a **true output y** and a **model-generated predicted output \hat{y}** for each data point.




Evaluation of Classification Models (1/3)

- ▶ In a **classification problem**, there exists a **true output y** and a **model-generated predicted output \hat{y}** for each data point.
- ▶ The results for each instance point can be assigned to one of **four categories**:
 - True Positive (TP)
 - True Negative (TN)
 - False Positive (FP)
 - False Negative (FN)



Evaluation of Classification Models (2/3)

- ▶ True Positive (TP): the label y is positive and prediction \hat{y} is also positive.
- ▶ True Negative (TN): the label y is negative and prediction \hat{y} is also negative.



Evaluation of Classification Models (3/3)

- ▶ False Positive (FP): the label y is negative but prediction \hat{y} is positive (type I error).
- ▶ False Negative (FN): the label y is positive but prediction \hat{y} is negative (type II error).





Why Pure Accuracy Is Not A Good Metric?

- ▶ **Accuracy:** how **close** the **prediction** is to the **true value**.



Why Pure Accuracy Is Not A Good Metric?

- ▶ **Accuracy:** how **close** the **prediction** is to the **true value**.
- ▶ Assume a highly **unbalanced dataset**
- ▶ E.g., a dataset where **95%** of the data points are **not fraud** and **5%** of the data points are **fraud**.



Why Pure Accuracy Is Not A Good Metric?

- ▶ **Accuracy:** how **close** the **prediction** is to the **true value**.
- ▶ Assume a highly **unbalanced dataset**
- ▶ E.g., a dataset where **95%** of the data points are **not fraud** and **5%** of the data points are **fraud**.
- ▶ A **naive classifier** that **predicts not fraud**, regardless of input, will be **95% accurate**.




Why Pure Accuracy Is Not A Good Metric?

- ▶ **Accuracy:** how **close** the **prediction** is to the **true value**.
- ▶ Assume a highly **unbalanced dataset**
- ▶ E.g., a dataset where **95%** of the data points are **not fraud** and **5%** of the data points are **fraud**.
- ▶ A **naive classifier** that **predicts not fraud**, regardless of input, will be **95% accurate**.
- ▶ For this reason, metrics like **precision** and **recall** are typically used.

Precision

- ▶ It is the **accuracy** of the **positive predictions**.


$$\text{Precision} = p(y = 1 \mid \hat{y} = 1) = \frac{\text{TP}}{\text{TP} + \text{FP}}$$



Recall

- ▶ Is the **ratio** of positive instances that are correctly detected by the classifier.
- ▶ Also called **sensitivity** or **true positive rate (TPR)**.

$$\text{Recall} = p(\hat{y} = 1 \mid y = 1) = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{Recall} = \frac{\text{Green Circle}}{\text{Green Square}}$$
A diagram illustrating the concept of recall. It consists of two overlapping shapes: a green circle at the top and a green square below it. The intersection of the circle and the square is also green, representing the true positives. The area of the intersection divided by the total area of the square represents the fraction of true positives relative to all actual positives (the union of the circle and the square).




F1 Score

- ▶ *F1 score*: combine precision and recall into a single metric.
- ▶ The harmonic mean of precision and recall.
- ▶ *F1* only gets high score if both recall and precision are high.


$$F1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}}$$

Confusion Matrix

- The **confusion matrix** is $K \times K$, where K is the **number of classes**.



Confusion Matrix - Example



$$TP = 3, TN = 5, FP = 1, FN = 2$$


$$\text{Precision} = \frac{TP}{TP + FP} = \frac{3}{3 + 1} = \frac{3}{4}$$

$$\text{Recall (TPR)} = \frac{TP}{TP + FN} = \frac{3}{3 + 2} = \frac{3}{5}$$

$$\text{FPR} = \frac{FP}{TN + FP} = \frac{1}{5 + 1} = \frac{1}{6}$$



Precision-Recall Tradeoff

- ▶ Precision-recall tradeoff: increasing precision reduces recall, and vice versa.



Precision-Recall Tradeoff

- ▶ Precision-recall tradeoff: increasing precision reduces recall, and vice versa.



The ROC Curve

- ▶ True positive rate (TPR) (recall): $p(\hat{y} = 1 \mid y = 1)$
- ▶ False positive rate (FPR): $p(\hat{y} = 1 \mid y = 0)$


$$\text{Recall} = \frac{\text{Green}}{\text{Total}}$$

$$\text{FPR} = \frac{\text{Red}}{\text{Total}}$$

The ROC Curve

- ▶ True positive rate (TPR) (recall): $p(\hat{y} = 1 \mid y = 1)$
- ▶ False positive rate (FPR): $p(\hat{y} = 1 \mid y = 0)$
- ▶ The **receiver operating characteristic (ROC)** curves summarize the **trade-off** between the TPR and FPR for a model using different probability **thresholds**.

$$\text{Recall} = \frac{\text{Green}}{\text{Total}}$$
$$\text{FPR} = \frac{\text{Red}}{\text{Total}}$$





Summary

Summary

- ▶ Linear regression model $\hat{y} = \mathbf{w}^T \mathbf{x}$
 - Learning parameters \mathbf{w}
 - Cost function $J(\mathbf{w})$
 - Learn parameters: normal equation, gradient descent (batch, stochastic, mini-batch)
- ▶ Generalization
 - Overfitting vs. underfitting
 - Bias vs. variance
 - Regularization: Lasso regression, Ridge regression, ElasticNet
- ▶ Hyperparameters and cross-validation

Summary

- ▶ Binomial logistic regression
 - $y \in \{0, 1\}$
 - Sigmoid function
 - Minimize the cross-entropy
- ▶ Multinomial logistic regression
 - $y \in \{1, 2, \dots, k\}$
 - Softmax function
 - Minimize the cross-entropy
- ▶ Performance measurements
 - TP, TF, FP, FN
 - Precision, recall, F1
 - Threshold and ROC



Questions?